

# FaceMaskSSDMobileNetV3\_FINAL\_HEROKU\_DEPLOYED

October 16, 2021

## 1 Start

```
[ ]: pip install -U git+https://github.com/albumentations-team/albumentations
```

```
Collecting git+https://github.com/albumentations-team/albumentations
  Cloning https://github.com/albumentations-team/albumentations to /tmp/pip-req-
  build-q5pwd9zy
  Running command git clone -q https://github.com/albumentations-
  team/albumentations /tmp/pip-req-build-q5pwd9zy
  Requirement already satisfied: numpy>=1.11.1 in /usr/local/lib/python3.7/dist-
  packages (from albumentations==1.1.0) (1.19.5)
  Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages
  (from albumentations==1.1.0) (1.4.1)
  Requirement already satisfied: scikit-image>=0.16.1 in
  /usr/local/lib/python3.7/dist-packages (from albumentations==1.1.0) (0.16.2)
  Requirement already satisfied: PyYAML in /usr/local/lib/python3.7/dist-packages
  (from albumentations==1.1.0) (3.13)
  Requirement already satisfied: qudida>=0.0.4 in /usr/local/lib/python3.7/dist-
  packages (from albumentations==1.1.0) (0.0.4)
  Requirement already satisfied: opencv-python>=4.1.1 in
  /usr/local/lib/python3.7/dist-packages (from albumentations==1.1.0) (4.1.2.30)
  Requirement already satisfied: typing-extensions in
  /usr/local/lib/python3.7/dist-packages (from
  qudida>=0.0.4->albumentations==1.1.0) (3.7.4.3)
  Requirement already satisfied: scikit-learn>=0.19.1 in
  /usr/local/lib/python3.7/dist-packages (from
  qudida>=0.0.4->albumentations==1.1.0) (0.22.2.post1)
  Requirement already satisfied: opencv-python-headless>=4.0.1 in
  /usr/local/lib/python3.7/dist-packages (from
  qudida>=0.0.4->albumentations==1.1.0) (4.5.3.56)
  Requirement already satisfied: PyWavelets>=0.4.0 in
  /usr/local/lib/python3.7/dist-packages (from scikit-
  image>=0.16.1->albumentations==1.1.0) (1.1.1)
  Requirement already satisfied: imageio>=2.3.0 in /usr/local/lib/python3.7/dist-
  packages (from scikit-image>=0.16.1->albumentations==1.1.0) (2.4.1)
  Requirement already satisfied: networkx>=2.0 in /usr/local/lib/python3.7/dist-
  packages (from scikit-image>=0.16.1->albumentations==1.1.0) (2.6.3)
  Requirement already satisfied: matplotlib!=3.0.0,>=2.0.0 in
```

```

/usr/local/lib/python3.7/dist-packages (from scikit-
image>=0.16.1->albumentations==1.1.0) (3.2.2)
Requirement already satisfied: pillow>=4.3.0 in /usr/local/lib/python3.7/dist-
packages (from scikit-image>=0.16.1->albumentations==1.1.0) (7.1.2)
Requirement already satisfied: python-dateutil>=2.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib!=3.0.0,>=2.0.0->scikit-
image>=0.16.1->albumentations==1.1.0) (2.8.2)
Requirement already satisfied: pyparsing!=2.0.4,!2.1.2,!2.1.6,>=2.0.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib!=3.0.0,>=2.0.0->scikit-
image>=0.16.1->albumentations==1.1.0) (2.4.7)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib!=3.0.0,>=2.0.0->scikit-
image>=0.16.1->albumentations==1.1.0) (1.3.2)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.7/dist-
packages (from matplotlib!=3.0.0,>=2.0.0->scikit-
image>=0.16.1->albumentations==1.1.0) (0.10.0)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages
(from cycycler>=0.10->matplotlib!=3.0.0,>=2.0.0->scikit-
image>=0.16.1->albumentations==1.1.0) (1.15.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-
packages (from scikit-learn>=0.19.1->qudida>=0.0.4->albumentations==1.1.0)
(1.0.1)

```

```
[ ]: from google.colab import drive
drive.mount("/content/gdrive")
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force\_remount=True).

```
[ ]: !nvidia-smi
```

Thu Oct 14 15:21:18 2021

```

+-----+
| NVIDIA-SMI 470.74      Driver Version: 460.32.03      CUDA Version: 11.2      |
+-----+-----+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                           MIG M. |
+-----+-----+-----+-----+
|   0   Tesla P100-PCIE...    Off   | 00000000:00:04.0 Off |                    0 |
| N/A   46C    P0      28W / 250W |      0MiB / 16280MiB |      0%      Default |
|                                           N/A |
+-----+-----+-----+-----+

+-----+
| Processes: |
| GPU   GI    CI          PID    Type    Process name                        GPU Memory |
|          ID    ID                                   Usage   |
+-----+

```

```
|=====|
| No running processes found |
+-----+
```

## 2 Modules

```
[ ]: # image
import cv2
import matplotlib.pyplot as plt
import seaborn as sns

# DL
import torch
import torchvision
from torch.utils.data import Dataset, DataLoader
from torch import nn as nn

# file
from glob import glob
import pandas as pd
import numpy as np
import os
import xml.etree.ElementTree as ET
from sklearn.model_selection import train_test_split
from tqdm import tqdm
import json

# augmentation
import albumentations as A
```

## 3 Data Extraction

```
[ ]: !mkdir csvFolder
```

### 3.1 Data 1

<https://www.kaggle.com/andrewmvd/face-mask-detection>

```
[ ]:
```

```
!curl --header 'Host: storage.googleapis.com' --user-agent 'Mozilla/5.0 (X11;
↳ Ubuntu; Linux x86_64; rv:93.0) Gecko/20100101 Firefox/93.0' --header 'Accept:
↳ text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/
↳ webp,*/*;q=0.8' --header 'Accept-Language: en-US,en;q=0.5' --referer 'https://
↳ /www.kaggle.com/' --header 'DNT: 1' --header 'Upgrade-Insecure-Requests: 1'
↳ --header 'Sec-Fetch-Dest: document' --header 'Sec-Fetch-Mode: navigate'
↳ --header 'Sec-Fetch-Site: cross-site' --header 'Sec-Fetch-User: ?1' 'https://
↳ storage.googleapis.com/kaggle-data-sets/667889/1176415/bundle/archive.zip?
↳ X-Goog-Algorithm=GOOG4-RSA-SHA256&X-Goog-Credential=gcp-kaggle-com%40kaggle-161607.
↳ iam.gserviceaccount.
↳ com%2F20211014%2Fauto%2Fstorage%2Fgoog4_request&X-Goog-Date=20211014T112923Z&X-Goog-Expires
↳ --output 'archive.zip'
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100	397M	100	397M	0	0	46.2M	0
0:00:08	0:00:08	--:--:--	48.9M				

```
[ ]: !unzip /content/archive.zip -d /content/data1
```

```
[ ]: !rm /content/archive.zip
```

```
[ ]: dataOneImagesPath = "/content/data1/images/"
dataOneAnnotationPath = "/content/data1/annotations/"
```

```
[ ]: import xml.etree.ElementTree as ET
```

```
#convert the object annotation from XML file to a dataframe
```

```
def xml_to_df(path):
```

```
    xml_list = []
```

```
    for xml_file in glob(path + '/*.xml'):
```

```
        tree = ET.parse(xml_file)
```

```
        root = tree.getroot()
```

```
        for member in root.findall('object'):
```

```
            ## chnage object index number, chnages with xml file
```

```
            xmin = float(member[5][0].text)
```

```
            ymin = float(member[5][1].text)
```

```
            xmax = float(member[5][2].text)
```

```
            ymax = float(member[5][3].text)
```

```
            img_width = float(root.find('size')[0].text)
```

```
            img_height = float(root.find('size')[1].text)
```

```
            aspectRatio = img_width / img_height
```

```

        box_width = (xmax-xmin)
        box_height = (ymax-ymin)
        bboxArea = box_width * box_height

        if xmax <= img_width and ymax <= img_height and bboxArea > 400:

            value = (dataOneImagesPath + root.find('filename').text,
                    member[0].text,
                    xmin,
                    ymin,
                    xmax,
                    ymax,
                    )

            xml_list.append(value)

        column_name = ['filename',
                       'class', 'xmin', 'ymin', 'xmax', 'ymax']

        xml_df = pd.DataFrame(xml_list, columns=column_name)
        return xml_df

df1 = xml_to_df(dataOneAnnotationPath)

df1.to_csv(os.path.join("/content/csvFolder", "data1.csv"), index=False)

```

```
[ ]: df1.shape
```

```
[ ]: (2346, 6)
```

```
[ ]: df1['class'].value_counts()
```

```
[ ]: with_mask          1926
      without_mask      337
      mask_wearred_incorrect  83
      Name: class, dtype: int64
```

## 3.2 Data 2

<https://www.kaggle.com/wobotintelligence/face-mask-detection-dataset>

```
[ ]:
```

```
!curl --header 'Host: storage.googleapis.com' --user-agent 'Mozilla/5.0 (X11;
↳ Ubuntu; Linux x86_64; rv:93.0) Gecko/20100101 Firefox/93.0' --header 'Accept:
↳ text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/
↳ webp,*/*;q=0.8' --header 'Accept-Language: en-US,en;q=0.5' --referer 'https:/
↳ /www.kaggle.com/' --header 'DNT: 1' --header 'Alt-Used: storage.googleapis.
↳ com' --header 'Upgrade-Insecure-Requests: 1' --header 'Sec-Fetch-Dest:
↳ document' --header 'Sec-Fetch-Mode: navigate' --header 'Sec-Fetch-Site:
↳ cross-site' --header 'Sec-Fetch-User: ?1' 'https://storage.googleapis.com/
↳ kaggle-data-sets/710024/1246711/bundle/archive.zip?
↳ X-Goog-Algorithm=GOOG4-RSA-SHA256&X-Goog-Credential=gcp-kaggle-com%40kaggle-161607.
↳ iam.gserviceaccount.
↳ com%2F20211014%2Fauto%2Fstorage%2Fgoog4_request&X-Goog-Date=20211014T112951Z&X-Goog-Expires
↳ --output 'archive.zip'
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100 2558M	100 2558M	0 0	54.2M 0	0:00:47	0:00:47	--:--:--	68.4M

```
[ ]: !unzip /content/archive.zip -d /content/data2
```

```
[ ]: !rm /content/archive.zip
```

```
[ ]: dataTwoImagesPath = "/content/data2/Medical mask/Medical mask/Medical Mask/
↳ images"
dataTwoAnnotationPath = "/content/data2/Medical mask/Medical mask/Medical Mask/
↳ annotations"
```

```
[ ]: def dataTwoJson2DF(path):

    rowList = []

    for file in glob(os.path.join(path, "*")):

        f = open(file)
        jsonDict = json.load(f)

        filename = os.path.join(dataTwoImagesPath, jsonDict['FileName'])
        noOfAnno = jsonDict['NumOfAnno']
        Annotations = jsonDict['Annotations']
        for i in range(noOfAnno):
            classname = Annotations[i]['classname']
            bbox = Annotations[i]['BoundingBox']
            xmin = float(bbox[0])
            ymin = float(bbox[1])
            xmax = float(bbox[2])
            ymax = float(bbox[3])
```

```

        value = (
            filename, classname, xmin, ymin, xmax, ymax
        )

        boxArea = (ymax-ymin) * (xmax-xmin)

        if boxArea > 400.0:
            rowList.append(value)

    column_name = ['filename', 'class', 'xmin', 'ymin', 'xmax', 'ymax']

    return pd.DataFrame(rowList, columns=column_name)

# convert data2 to df
df2 = dataTwoJson2DF(dataTwoAnnotationPath)

```

```
[ ]: df2.shape
```

```
[ ]: (15178, 6)
```

```

[ ]: df2["class"].replace({"face_with_mask_incorrect": "mask_wearred_incorrect",
                           "face_with_mask": "with_mask",
                           "mask_surgical": "removeCat",
                           "mask_colorful": "removeCat",
                           "face_no_mask": "without_mask",
                           "gas_mask": "removeCat",

                           "balaclava_ski_mask": "removeCat",
                           "turban": "removeCat",
                           "hood": "removeCat",
                           "other": "removeCat",
                           "goggles": "removeCat",
                           "scarf_bandana": "removeCat",
                           "hair_net": "removeCat",
                           "sunglasses": "removeCat",
                           "hat": "removeCat",
                           "face_other_covering": "removeCat",
                           "hijab_niqab": "removeCat",
                           "face_shield": "removeCat",
                           "eyeglasses": "removeCat",
                           "helmet": "removeCat",

                           },

```

```
inplace=True)
```

```
[ ]: df2 = df2[df2['class'] != "removeCat"]
```

```
[ ]: df2.to_csv("/content/csvFolder/data2.csv", index=False)
```

```
[ ]: df2['class'].value_counts()
```

```
[ ]: with_mask          4167
      without_mask      1562
      mask_wearred_incorrect  148
      Name: class, dtype: int64
```

## 4 Merge Dataset

```
[ ]: df1['class'].value_counts()
```

```
[ ]: with_mask          1926
      without_mask      337
      mask_wearred_incorrect  83
      Name: class, dtype: int64
```

```
[ ]: df2['class'].value_counts()
```

```
[ ]: with_mask          4167
      without_mask      1562
      mask_wearred_incorrect  148
      Name: class, dtype: int64
```

```
[ ]: DF = df1.append(df2, ignore_index=True)
```

```
[ ]: DF.shape
```

```
[ ]: (8223, 6)
```

```
[ ]: DF['class'].value_counts()
```

```
[ ]: with_mask          6093
      without_mask      1899
      mask_wearred_incorrect  231
      Name: class, dtype: int64
```



## 5 Import Scripts

```
[ ]: !git clone https://github.com/pytorch/vision.git
```

```
Cloning into 'vision'...
remote: Enumerating objects: 42731, done.
remote: Counting objects: 100% (8207/8207), done.
remote: Compressing objects: 100% (2126/2126), done.
remote: Total 42731 (delta 6356), reused 7562 (delta 5921), pack-reused 34524
Receiving objects: 100% (42731/42731), 65.11 MiB | 28.59 MiB/s, done.
Resolving deltas: 100% (32907/32907), done.
```

```
[ ]: # import a python file as a module in Notebook
```

```
!ln -s ./vision/references/detection/engine.py engine.py
!ln -s ./vision/references/detection/train.py train.py
!ln -s ./vision/references/detection/utils.py utils.py
!ln -s ./vision/references/detection/coco_utils.py coco_utils.py
!ln -s ./vision/references/detection/group_by_aspect_ratio.py↵
↪group_by_aspect_ratio.py
!ln -s ./vision/references/detection/presets.py presets.py
!ln -s ./vision/references/detection/coco_eval.py coco_eval.py
!ln -s ./vision/references/detection/transforms.py transforms.py
```

```
[ ]: # to use pyprotocols
!pip install -q pycocotools
```

## 6 Train Eval Split

```
[ ]: from sklearn.model_selection import train_test_split
```

```
[ ]: filenames = DF['filename'].unique().tolist()

TRAIN_IMAGES, VALID_IMAGES = train_test_split(filenames, test_size=0.2)

len(TRAIN_IMAGES), len(VALID_IMAGES)
```

```
[ ]: (3406, 852)
```

```
[ ]: def get_dataframe_from_Images(DF, images):
    for i, image in enumerate(images):
        if i == 0:
            first_df = DF[DF['filename']==image]
            return_df = first_df
        else:
            temp_df = DF[DF['filename']==image]
```

```

        return_df = return_df.append(temp_df, ignore_index=True)

    return return_df

TRAIN_DF = get_dataframe_from_Images(DF, TRAIN_IMAGES)
VALID_DF = get_dataframe_from_Images(DF, VALID_IMAGES)

```

```
[ ]: TRAIN_DF['filename'].unique().shape, VALID_DF['filename'].unique().shape
```

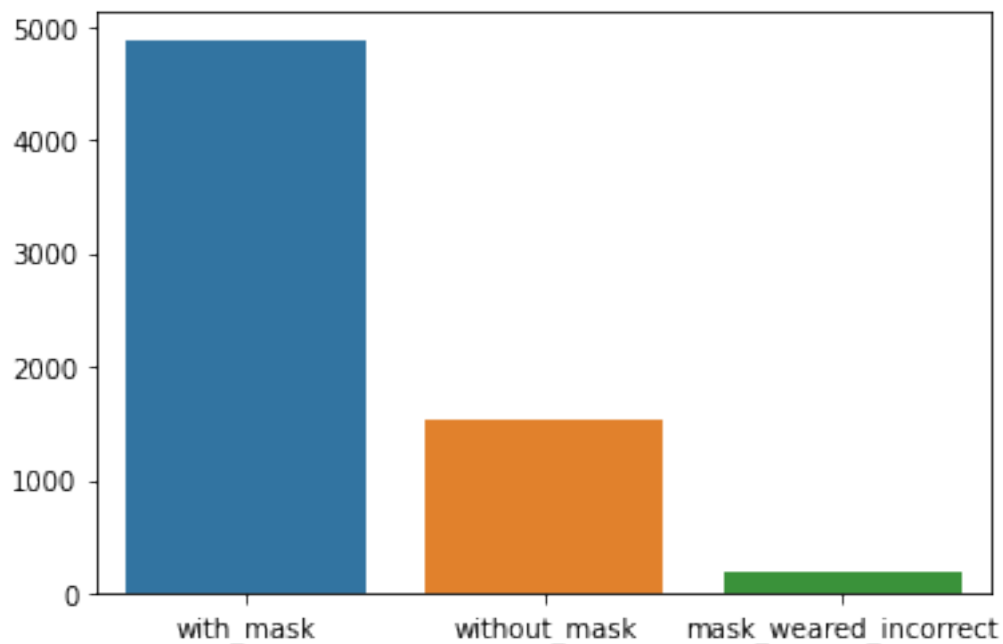
```
[ ]: ((3406,), (852,))
```

```
[ ]: sns.barplot(TRAIN_DF['class'].value_counts().index, TRAIN_DF['class'].
    ↳value_counts().values)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f17d44cc3d0>
```




```
[ ]: sns.barplot(VALID_DF['class'].value_counts().index, VALID_DF['class'].
    ↳value_counts().values)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning:

Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f17d340ab10>
```



## 7 Visualize Dataset

```
[ ]: from albumentations.pytorch.transforms import ToTensorV2
```

```
[ ]: # A.Resize(512, 512, always_apply=True)

def get_aug_pipeline():

    transforms = []
    # transforms.append( )
    transforms.append(ToTensorV2(p=1))

    return A.Compose(
        transforms, bbox_params=A.BboxParams(format='pascal_voc',
        ↪label_fields=['labels']),
    )
```

```
[ ]: from matplotlib import pyplot as plt
import matplotlib.patches as patches

def show_output_with_bbox(filename, bboxes, labels, transform):

    # image = plt.imread(filename)

    image = cv2.imread(filename, cv2.IMREAD_ANYCOLOR) # cv2.IMREAD_COLOR
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB).astype(np.float32)
    image = image/255.0

    aug_pipeline = get_aug_pipeline()
    transformed = aug_pipeline(image=image, bboxes=bboxes, labels=labels)

    img = transformed['image']
    bboxes = torch.as_tensor(transformed['bboxes'])
    bboxes = bboxes.detach().numpy()
    labels = transformed['labels']

    print(img.shape)

    img_height = img.shape[1]
    img_width = img.shape[2]

    fig, ax = plt.subplots(figsize=(5,5))

    ax.imshow(img.permute(1,2,0).numpy())

    for bbox, class_name in zip(bboxes, labels):
        xmin = bbox[0]
        ymin = bbox [1]
        width = bbox[2] - xmin
        height = bbox[3] - ymin

        # xmin, ymin, width, height = [x for x in [(xcenter-(width/2)),
        ↪(ycenter-(height/2)), width, height]]
        rect = patches.Rectangle((xmin, ymin), width, height, linewidth=1,
        ↪edgecolor='r', facecolor='none')
        ax.add_patch(rect)
        ax.text(xmin, ymin, class_name, color='w')

    plt.show()
```

```
#####

trainFiles = TRAIN_DF['filename'].unique().tolist()
trainFiles = trainFiles[:50]

for filename in trainFiles:
    records = TRAIN_DF[TRAIN_DF['filename']==filename]

    bboxes = records[['xmin', 'ymin', 'xmax', 'ymax']].values
    labels = records[['class']].values
    labels = [x[0] for x in labels]
    area = (bboxes[:,2]-bboxes[:,0]) * (bboxes[:,3]-bboxes[:,1])
    print(area)

    show_output_with_bbox(filename, bboxes, labels, get_aug_pipeline())
```

[ ]: *# Image Verification*

```
trainFiles = TRAIN_DF['filename'].unique().tolist()

for filename in trainFiles:

    if filename == "/content/data1/images/maksssksksss719.png":

        records = TRAIN_DF[TRAIN_DF['filename']==filename]

        bboxes = records[['xmin', 'ymin', 'xmax', 'ymax']].values
        labels = records[['class']].values
        labels = [x[0] for x in labels]
        area = (bboxes[:,2]-bboxes[:,0]) * (bboxes[:,3]-bboxes[:,1])
        print(area)

        show_output_with_bbox(filename, bboxes, labels, get_aug_pipeline())

    break
```

```
[ 616.  540.  900. 1512.]
torch.Size([3, 226, 400])
```



## 8 Dataset and Dataloader

```
[ ]: import random
# random.seed(7)

[ ]: from albumentations.pytorch.transforms import ToTensorV2

# https://albumentations.ai/docs/getting_started/bounding_boxes_augmentation/
# https://albumentations.ai/docs/api_reference/core/composition/

global class_int_dict
class_int_dict = {
    "background":0,
    "with_mask":1,
    "without_mask":2,
    "mask_wearred_incorrect":3
}

[ ]: def get_geometric_aug(aug):
    return A.Compose([A.Resize(320, 320, p=1.0),
                      aug,
                      A.Resize(320, 320, p=1.0),
                      ToTensorV2(p=1)],
                    bbox_params=A.BboxParams(format='pascal_voc', min_area=0.
↪0, min_visibility=0.0, label_fields=['labels']))

def get_noise_aug(aug):
    return A.Compose([A.Resize(320, 320, p=1.0),
```

```

        aug,
        ToTensorV2(p=1)],
        bbox_params=A.BboxParams(format='pascal_voc', min_area=0.
→0, min_visibility=0.0, label_fields=['labels']))

def get_resize_aug():
    return A.Compose([A.Resize(320, 320, p=1.0),
        ToTensorV2(p=1)],
        bbox_params=A.BboxParams(format='pascal_voc', min_area=0.
→0, min_visibility=0.0, label_fields=['labels']))

```

```

[ ]: class FacemaskDataset(Dataset):

    def __init__(self, DF, train):

        self.train = train
        self.DF = DF
        self.filenamees = self.DF['filename'].unique()

        self.resize = get_resize_aug()
        self.rotate = get_geometric_aug(A.Rotate(limit = 10, p=1.0))
        self.transpose = get_geometric_aug(A.Transpose(p=1.0))
        self.h_flip = get_geometric_aug(A.HorizontalFlip(p=1.0))
        self.rand_flip = get_geometric_aug(A.Flip(p=1.0))
        self.center_crop_400 = get_geometric_aug(A.CenterCrop(height=400,
→width=400, p=1.0))
        self.crop_pad_zoom_in_10 = get_geometric_aug(A.CropAndPad(percent=-0.
→08, p=1))
        self.crop_pad_zoom_out_10 = get_geometric_aug(A.CropAndPad(percent=0.
→10, p=1))

        self.gaussian_blur = get_noise_aug(A.GaussianBlur(blur_limit=(5, 7),
→sigma_limit=0, always_apply=True, p=1))
        self.med_blur = get_noise_aug(A.MedianBlur(blur_limit=5,
→always_apply=True, p=1))
        self.glass_blur = get_noise_aug(A.GlassBlur(p=1))
        self.mult_noise = get_noise_aug(A.MultiplicativeNoise(multiplier=(1.1,
→2.1), p=1))
        self.gaussian_noise = get_noise_aug(A.GaussNoise(p=1))
        self.gray = get_noise_aug(A.ToGray(p=1))
        self.rand_contrast = get_noise_aug(A.
→RandomBrightnessContrast(brightness_limit=0.2, contrast_limit=0.7, p=1))
        self.aa_sharpen = get_noise_aug(A.Sharpen(alpha=(0.1, 1), lightness=(0.
→8, 1.0), p=1.0))

        self.transforms = None

```

```

def __len__(self):
    return len(self.filenamees)

def load_images_boxes(self, index):

    image_id = self.filenamees[index]

    records = self.DF[self.DF['filename'] == image_id]
    img = cv2.imread(image_id, cv2.IMREAD_ANYCOLOR) # cv2.IMREAD_COLOR
    img = cv2.cvtColor(img, cv2.COLOR_BGRA2RGB).astype(np.float32)
    img = img/255.0

    bboxes = records[['xmin', 'ymin', 'xmax', 'ymax']].values.tolist()
    class_labels = records[['class']].values.reshape(1, -1)[0]
    class_labels = [class_int_dict[i] for i in class_labels]

    return img, bboxes, class_labels

def __getitem__(self, index):

    img, bboxes, class_labels = self.load_images_boxes(index)

    rand_number = np.round(random.random(), 3)

    # print(f"random number: {rand_number}")

    if self.train:

        if rand_number >= 0.65:
            # print("iaa_sharpen")
            self.transforms = self.aa_sharpen
            transformed = self.transforms(image=img, bboxes=bboxes,
→labels=class_labels)

        elif rand_number >= 0.52:
            # print("gray")
            self.transforms = self.gray
            transformed = self.transforms(image=img, bboxes=bboxes,
→labels=class_labels)

        elif rand_number >= 0.39:
            # print("h_flip")
            self.transforms = self.h_flip

```



```

        transformed = self.transforms(image=img, bboxes=bboxes,
→labels=class_labels)

        elif rand_number >= 0.26:
            # print("gaussian_blur")
            self.transforms = self.gaussian_blur
            transformed = self.transforms(image=img, bboxes=bboxes,
→labels=class_labels)

        elif rand_number >= 0.13:
            # print("rotate")
            self.transforms = self.rotate
            transformed = self.transforms(image=img, bboxes=bboxes,
→labels=class_labels)

        elif rand_number >= 0.0:
            # print("med_blur")
            self.transforms = self.med_blur
            transformed = self.transforms(image=img, bboxes=bboxes,
→labels=class_labels)

        else:
            self.transforms = self.resize
            transformed = self.transforms(image=img, bboxes=bboxes,
→labels=class_labels)

    img = transformed['image']
    bboxes = torch.as_tensor(transformed['bboxes'], dtype=torch.float32)

    num_objs = len(bboxes)
    iscrowd = torch.zeros((num_objs,), dtype=torch.int64)
    class_label = torch.tensor(transformed['labels'])
    iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

    target = {}
    target["labels"] = class_label # tensor
    target["boxes"] = bboxes # tensor
    target['image_id'] = torch.tensor([index]) # string
    target['iscrowd'] = iscrowd
    target['area'] = (bboxes[:, 3] - bboxes[:, 1]) * (bboxes[:, 2] -
→bboxes[:, 0])

    return img, target

```

```
[ ]: train_dataset = FacemaskDataset(DF=TRAIN_DF, train=True)

valid_dataset = FacemaskDataset(DF=VALID_DF, train=False)

def collate_fn(batch):
    return tuple(zip(*batch))

train_data_loader = DataLoader(
    train_dataset,
    batch_size = 30,
    shuffle = True,
    num_workers = 2,
    collate_fn = collate_fn
)

valid_data_loader = DataLoader(
    valid_dataset,
    batch_size = 20,
    shuffle = True,
    num_workers = 2,
    collate_fn = collate_fn
)
```

```
[ ]: for i in range(len(train_data_loader)):
    train_dataset[i]
```

```
[ ]: from matplotlib import pyplot as plt
import matplotlib.patches as patches

class_dict = {
    0:"background",
    1:"with_mask",
    2:"without_mask",
    3:"mask_wearred_incorrect"
}

def show_output_with_bbox(img_tensor, bbox_tensor, labels_tensor):
    img = img_tensor.permute(1,2,0).detach().cpu().numpy()
    bboxes = bbox_tensor.detach().cpu().numpy()
    labels = labels_tensor.detach().cpu().numpy()

    fig, ax = plt.subplots(figsize=(5,5))
```

```

ax.imshow(img)

for bbox, label in zip(bboxes, labels):

    class_name = class_dict[label]
    xmin = bbox[0]
    ymin = bbox [1]
    width = bbox[2] - xmin
    height = bbox[3] - ymin

    rect = patches.Rectangle((xmin, ymin), width, height, linewidth=1,
↪edgecolor='r', facecolor='none')
    ax.add_patch(rect)
    ax.text(xmin, ymin, class_name, color='w')

    # print(f"class: {class_name}, area: {height*width}")

plt.show()

#####

for i in range(TRAIN_DF.shape[0]):
    image, target = train_dataset[i]
    bboxes = target['boxes']
    labels = target['labels']
    print(target['area'])

    show_output_with_bbox(image, bboxes, labels)

    if i == 10:
        break

```

```

tensor([ 5128.9180, 10574.6387,  3902.8801,  2738.6248])

```




Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
tensor([10915.2666, 3541.3977])
```

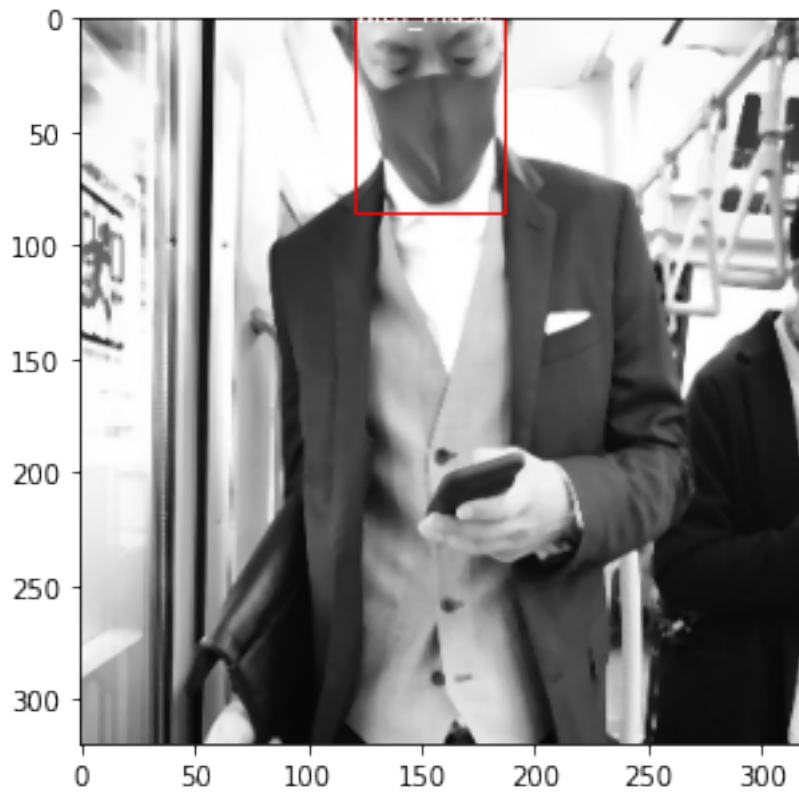


Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).


```
tensor([7242.7974])
```



tensor([5629.1323])

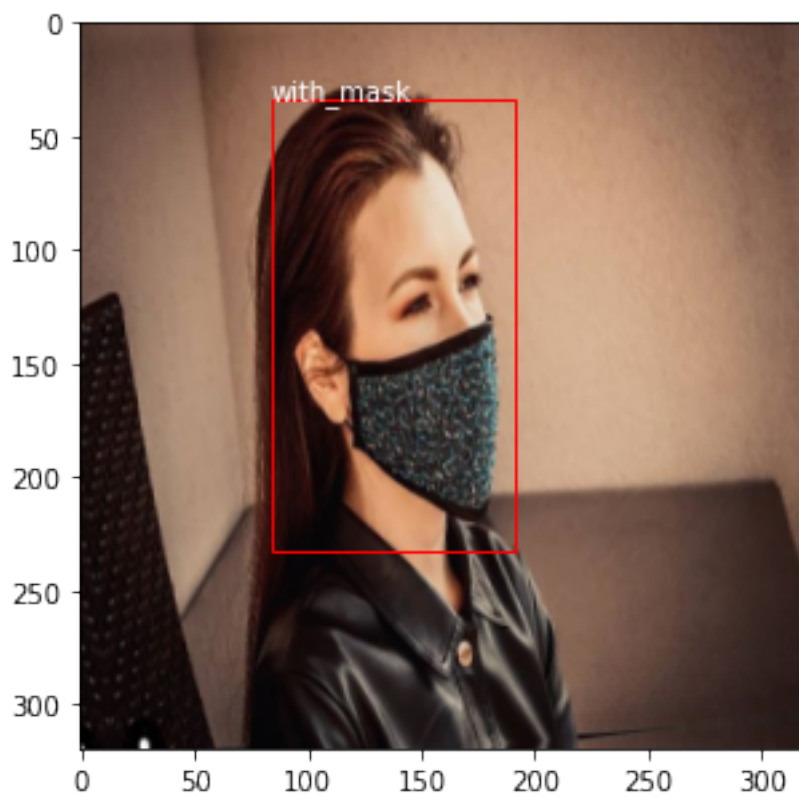


tensor([54700.5156])



tensor([21445.1914])

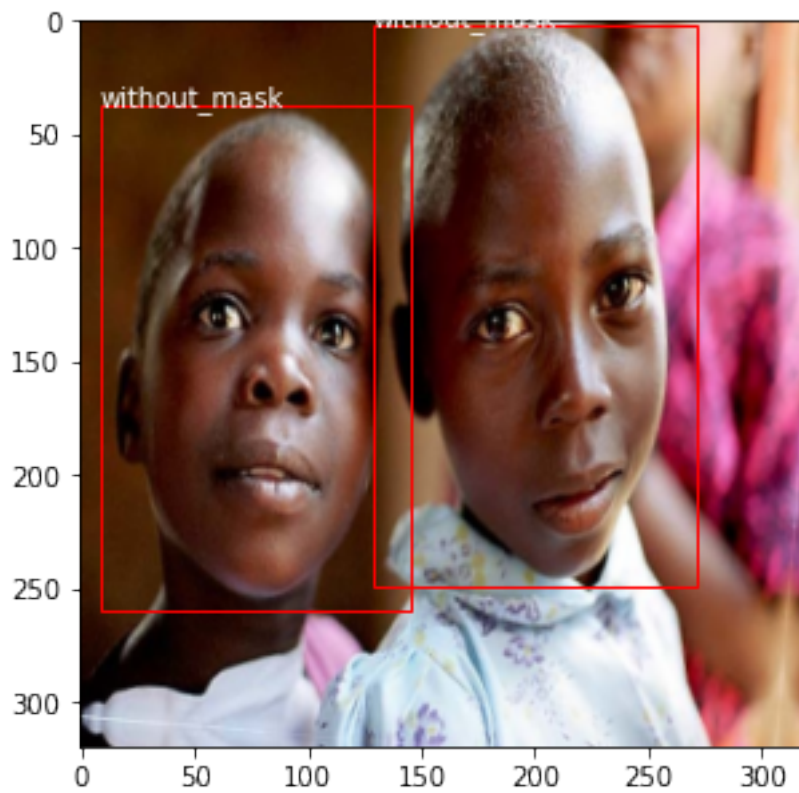




tensor([1377.4155, 475.7522])




tensor([35510.5898, 30364.7734])




Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
tensor([7104.6934])
```

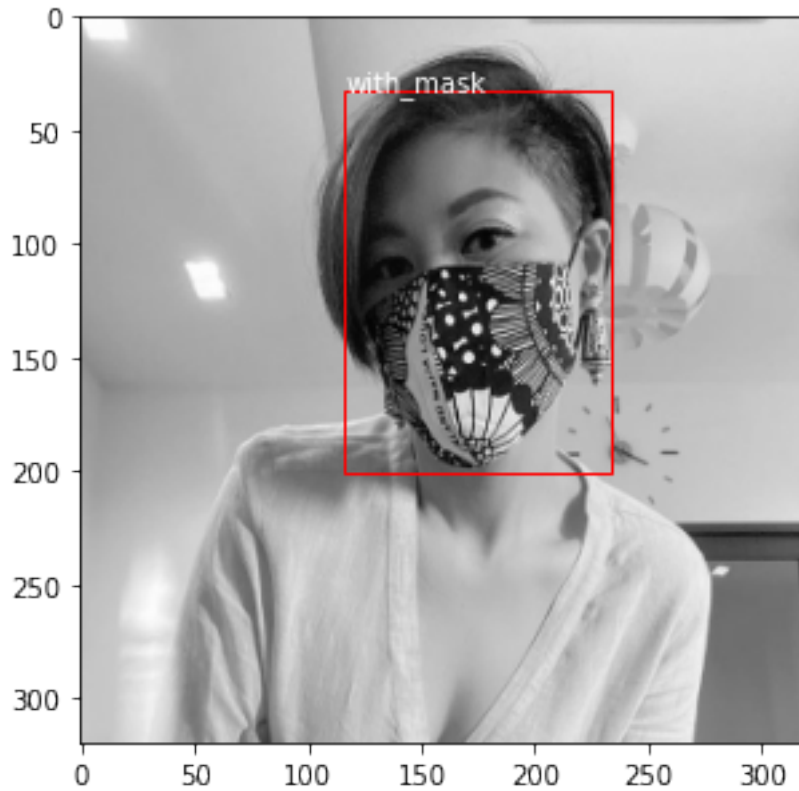


Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
tensor([6751.7690])
```



tensor([19875.0781])



```
[ ]: !nvidia-smi
```

Thu Oct 14 15:24:02 2021

```
+-----+
| NVIDIA-SMI 470.74          Driver Version: 460.32.03   CUDA Version: 11.2   |
+-----+-----+
| GPU  Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                      | MIG M. |
+=====+=====+
|   0   Tesla P100-PCIE...    Off  | 00000000:00:04.0 Off |                    0 |
| N/A   46C    P0      28W / 250W |      0MiB / 16280MiB |      0%      Default |
|                                       |                      | N/A |
+-----+-----+

```

```
+-----+
| Processes:
| GPU   GI    CI          PID    Type    Process name                        GPU Memory
|       ID    ID                                   Usage
+=====+
| No running processes found
+-----+

```

## 9 Device

```
[ ]: def getDevice():  
      return torch.device("cuda") if torch.cuda.is_available() else torch.  
      ↪device("cpu")
```

## 10 Weighted Loss SSD

```
[ ]: TRAIN_DF['class'].value_counts()
```

```
[ ]: with_mask          5525  
      without_mask      1718  
      mask_wearred_incorrect  206  
      Name: class, dtype: int64
```

```
[ ]: total = 5525 + 1718 + 206  
      total
```

```
[ ]: 7449
```

```
[ ]:
```

```
[ ]: global classWeight  
  
      classWeight = torch.tensor([1, 1.3, 4.3, 36.16], dtype=torch.float32, ↪  
      ↪device=getDevice())
```

```
[ ]: from typing import Any, Dict, List, Optional, Tuple  
      from torch import nn  
      import torch.nn.functional as F  
      from collections import OrderedDict  
      import torch  
      import torchvision  
      import torchvision.models.detection._utils as det_utils  
      from torch import nn, Tensor  
      from torchvision.ops import boxes as box_ops
```

```
[ ]: # Weighted SSD Implementation  
  
      class WeightedLossSSD(nn.Module):  
          """  
          Implements SSD architecture from `"SSD: Single Shot MultiBox Detector" ↪  
          ↪<https://arxiv.org/abs/1512.02325>`_.
```



The input to the model is expected to be a list of tensors, each of shape  $\rightarrow [C, H, W]$ , one for each image, and should be in 0-1 range. Different images can have different  $\rightarrow$  sizes but they will be resized to a fixed size before passing it to the backbone.

The behavior of the model changes depending if it is in training or  $\rightarrow$  evaluation mode.

During training, the model expects both the input tensors, as well as a  $\rightarrow$  targets (list of dictionary), containing:

- boxes (`FloatTensor[N, 4]`): the ground-truth boxes in `[x1, y1,  $\rightarrow$  x2, y2]` format, with `0 <= x1 < x2 <= W` and `0 <= y1 < y2 <= H`.
- labels (`Int64Tensor[N]`): the class label for each ground-truth box

The model returns a `Dict[Tensor]` during training, containing the  $\rightarrow$  classification and regression losses.

During inference, the model requires only the input tensors, and returns  $\rightarrow$  the post-processed predictions as a `List[Dict[Tensor]]`, one for each input image. The fields  $\rightarrow$  of the Dict are as follows, where `N` is the number of detections:

- boxes (`FloatTensor[N, 4]`): the predicted boxes in `[x1, y1, x2,  $\rightarrow$  y2]` format, with `0 <= x1 < x2 <= W` and `0 <= y1 < y2 <= H`.
- labels (`Int64Tensor[N]`): the predicted labels for each detection
- scores (`Tensor[N]`): the scores for each detection

Args:

backbone (`nn.Module`): the network used to compute the features for the  $\rightarrow$  model.

It should contain an `out_channels` attribute with the list of the  $\rightarrow$  output channels of each feature map. The backbone should return a single Tensor or an  $\rightarrow$  `OrderedDict[Tensor]`.

anchor\_generator (`DefaultBoxGenerator`): module that generates the  $\rightarrow$  default boxes for a set of feature maps.

size (`Tuple[int, int]`): the width and height to which images will be  $\rightarrow$  rescaled before feeding them to the backbone.



```

        num_classes (int): number of output classes of the model (excluding the
        ↳background).
        image_mean (Tuple[float, float, float]): mean values used for input
        ↳normalization.
            They are generally the mean values of the dataset on which the
        ↳backbone has been trained
            on
        image_std (Tuple[float, float, float]): std values used for input
        ↳normalization.
            They are generally the std values of the dataset on which the
        ↳backbone has been trained on
        head (nn.Module, optional): Module run on top of the backbone features.
        ↳Defaults to a module containing
            a classification and regression module.
        score_thresh (float): Score threshold used for postprocessing the
        ↳detections.
        nms_thresh (float): NMS threshold used for postprocessing the
        ↳detections.
        detections_per_img (int): Number of best detections to keep after NMS.
        iou_thresh (float): minimum IoU between the anchor and the GT box so
        ↳that they can be
            considered as positive during training.
        topk_candidates (int): Number of best detections to keep before NMS.
        positive_fraction (float): a number between 0 and 1 which indicates the
        ↳proportion of positive
            proposals used during the training of the classification head. It
        ↳is used to estimate the negative to
            positive ratio.
        """
        __annotations__ = {
            'box_coder': det_utils.BoxCoder,
            'proposal_matcher': det_utils.Matcher,
        }

        def __init__(self, backbone: nn.Module, anchor_generator: torchvision.
        ↳models.detection.anchor_utils.DefaultBoxGenerator,
            size: Tuple[int, int], num_classes: int,
            image_mean: Optional[List[float]] = None, image_std:
        ↳Optional[List[float]] = None,
            head: Optional[nn.Module] = None,
            score_thresh: float = 0.01,
            nms_thresh: float = 0.45,
            detections_per_img: int = 200,
            iou_thresh: float = 0.5,
            topk_candidates: int = 400,
            positive_fraction: float = 0.25):

```

```

super().__init__()

self.backbone = backbone

self.anchor_generator = anchor_generator

self.box_coder = torchvision.models.detection._utils.
↳BoxCoder(weights=(10., 10., 5., 5.))

if head is None:
    if hasattr(backbone, 'out_channels'):
        out_channels = backbone.out_channels
    else:
        out_channels = det_utils.retrieve_out_channels(backbone, size)

    assert len(out_channels) == len(anchor_generator.aspect_ratios)

    num_anchors = self.anchor_generator.num_anchors_per_location()
    head = torchvision.models.detection.ssd.SSDHead(out_channels,
↳num_anchors, num_classes)
    self.head = head

self.proposal_matcher = det_utils.SSDMatcher(iou_thresh)

if image_mean is None:
    image_mean = [0.485, 0.456, 0.406]
if image_std is None:
    image_std = [0.229, 0.224, 0.225]
self.transform = torchvision.models.detection.transform.
↳GeneralizedRCNNTransform(min(size), max(size), image_mean, image_std,
                           size_divisible=1,
↳fixed_size=size)

self.score_thresh = score_thresh
self.nms_thresh = nms_thresh
self.detections_per_img = detections_per_img
self.topk_candidates = topk_candidates
self.neg_to_pos_ratio = (1.0 - positive_fraction) / positive_fraction

# used only on torchscript mode
self._has_warned = False

@torch.jit.unused
def eager_outputs(self, losses: Dict[str, Tensor],
                 detections: List[Dict[str, Tensor]]) -> Tuple[Dict[str,
↳Tensor], List[Dict[str, Tensor]]]:
    if self.training:

```

```

        return losses

    return detections

    def compute_loss(self, targets: List[Dict[str, Tensor]], head_outputs: Dict[str, Tensor], anchors: List[Tensor], matched_idxs: List[Tensor]) -> Dict[str, Tensor]:
        bbox_regression = head_outputs['bbox_regression']
        cls_logits = head_outputs['cls_logits']

        # Match original targets with default boxes
        num_foreground = 0
        bbox_loss = []
        cls_targets = []
        for (targets_per_image, bbox_regression_per_image, cls_logits_per_image, anchors_per_image, matched_idxs_per_image) in zip(targets, bbox_regression, cls_logits, anchors, matched_idxs):
            # produce the matching between boxes and targets
            foreground_idxs_per_image = torch.where(matched_idxs_per_image >= 0)[0]

            foreground_matched_idxs_per_image = matched_idxs_per_image[foreground_idxs_per_image]
            num_foreground += foreground_matched_idxs_per_image.numel()

            # Calculate regression loss
            matched_gt_boxes_per_image = targets_per_image['boxes'][foreground_matched_idxs_per_image]
            bbox_regression_per_image = bbox_regression_per_image[foreground_idxs_per_image, :]
            anchors_per_image = anchors_per_image[foreground_idxs_per_image, :]
            target_regression = self.box_coder.encode_single(matched_gt_boxes_per_image, anchors_per_image)
            bbox_loss.append(torch.nn.functional.smooth_l1_loss(bbox_regression_per_image, target_regression, reduction='sum'))

            # Estimate ground truth for class targets
            gt_classes_target = torch.zeros((cls_logits_per_image.size(0), ), dtype=targets_per_image['labels'].dtype, device=targets_per_image['labels'].device)
            gt_classes_target[foreground_idxs_per_image] = targets_per_image['labels'][foreground_matched_idxs_per_image]

```

```

        cls_targets.append(gt_classes_target)

    bbox_loss = torch.stack(bbox_loss)
    cls_targets = torch.stack(cls_targets)

    # Calculate classification loss
    num_classes = cls_logits.size(-1)
    cls_loss = F.cross_entropy(
        cls_logits.view(-1, num_classes),
        cls_targets.view(-1),
        weight=classWeight,
        reduction='none'
    ).view(cls_targets.size())

    # Hard Negative Sampling
    foreground_idx = cls_targets > 0
    num_negative = self.neg_to_pos_ratio * foreground_idx.sum(1,
→keepdim=True)
    # num_negative[num_negative < self.neg_to_pos_ratio] = self.
→neg_to_pos_ratio
    negative_loss = cls_loss.clone()
    negative_loss[foreground_idx] = -float('inf') # use -inf to detect
→positive values that crept in the sample
    values, idx = negative_loss.sort(1, descending=True)
    # background_idx = torch.logical_and(idx.sort(1)[1] < num_negative,
→torch.isfinite(values))
    background_idx = idx.sort(1)[1] < num_negative

    N = max(1, num_foreground)
    return {
        'bbox_regression': bbox_loss.sum() / N,
        'classification': (cls_loss[foreground_idx].sum() +
→cls_loss[background_idx].sum()) / N,
    }

    def forward(self, images: List[Tensor],
                targets: Optional[List[Dict[str, Tensor]]] = None) ->
→Tuple[Dict[str, Tensor], List[Dict[str, Tensor]]]:
        if self.training and targets is None:
            raise ValueError("In training mode, targets should be passed")

        if self.training:
            assert targets is not None
            for target in targets:
                boxes = target["boxes"]
                if isinstance(boxes, torch.Tensor):

```

```

        if len(boxes.shape) != 2 or boxes.shape[-1] != 4:
            raise ValueError("Expected target boxes to be a tensor"
                              "of shape [N, 4], got {}".format(
                                  boxes.shape))
        else:
            raise ValueError("Expected target boxes to be of type "
                              "Tensor, got {}".format(type(boxes)))

    # get the original image sizes
    original_image_sizes: List[Tuple[int, int]] = []
    for img in images:
        val = img.shape[-2:]
        assert len(val) == 2
        original_image_sizes.append((val[0], val[1]))

    # transform the input
    images, targets = self.transform(images, targets)

    # Check for degenerate boxes
    if targets is not None:
        for target_idx, target in enumerate(targets):
            boxes = target["boxes"]
            degenerate_boxes = boxes[:, 2:] <= boxes[:, :2]
            if degenerate_boxes.any():
                bb_idx = torch.where(degenerate_boxes.any(dim=1))[0][0]
                degen_bb: List[float] = boxes[bb_idx].tolist()
                raise ValueError("All bounding boxes should have positive_
→height and width."
                                " Found invalid box {} for target at index_
→{}."
                                .format(degen_bb, target_idx))

    # get the features from the backbone
    features = self.backbone(images.tensors)
    if isinstance(features, torch.Tensor):
        features = OrderedDict([('0', features)])

    features = list(features.values())

    # compute the ssd heads outputs using the features
    head_outputs = self.head(features)

    # create the set of anchors
    anchors = self.anchor_generator(images, features)

    losses = {}
    detections: List[Dict[str, Tensor]] = []

```

```

        if self.training:
            assert targets is not None

            matched_idxs = []
            for anchors_per_image, targets_per_image in zip(anchors, targets):
                if targets_per_image['boxes'].numel() == 0:
                    matched_idxs.append(torch.full((anchors_per_image.
→size(0),), -1, dtype=torch.int64,
                                                    device=anchors_per_image.
→device))
                    continue

                match_quality_matrix = box_ops.
→box_iou(targets_per_image['boxes'], anchors_per_image)
                matched_idxs.append(self.proposal_matcher(match_quality_matrix))

            losses = self.compute_loss(targets, head_outputs, anchors,
→matched_idxs)
            else:
                detections = self.postprocess_detections(head_outputs, anchors,
→images.image_sizes)
                detections = self.transform.postprocess(detections, images.
→image_sizes, original_image_sizes)

            if torch.jit.is_scripting():
                if not self._has_warned:
                    warnings.warn("SSD always returns a (Losses, Detections) tuple,
→in scripting")
                    self._has_warned = True
                return losses, detections
            return self.eager_outputs(losses, detections)

        def postprocess_detections(self, head_outputs: Dict[str, Tensor],
→image_anchors: List[Tensor],
                                image_shapes: List[Tuple[int, int]]) ->
→List[Dict[str, Tensor]]:
            bbox_regression = head_outputs['bbox_regression']
            pred_scores = F.softmax(head_outputs['cls_logits'], dim=-1)

            num_classes = pred_scores.size(-1)
            device = pred_scores.device

            detections: List[Dict[str, Tensor]] = []

            for boxes, scores, anchors, image_shape in zip(bbox_regression,
→pred_scores, image_anchors, image_shapes):

```

```

boxes = self.box_coder.decode_single(boxes, anchors)
boxes = box_ops.clip_boxes_to_image(boxes, image_shape)

image_boxes = []
image_scores = []
image_labels = []
for label in range(1, num_classes):
    score = scores[:, label]

    keep_idxs = score > self.score_thresh
    score = score[keep_idxs]
    box = boxes[keep_idxs]

    # keep only topk scoring predictions
    num_topk = min(self.topk_candidates, score.size(0))
    score, idxs = score.topk(num_topk)
    box = box[idxs]

    image_boxes.append(box)
    image_scores.append(score)
    image_labels.append(torch.full_like(score, fill_value=label,
↳dtype=torch.int64, device=device))

image_boxes = torch.cat(image_boxes, dim=0)
image_scores = torch.cat(image_scores, dim=0)
image_labels = torch.cat(image_labels, dim=0)

# non-maximum suppression
keep = box_ops.batched_nms(image_boxes, image_scores, image_labels,
↳self.nms_thresh)
keep = keep[:self.detections_per_img]

detections.append({
    'boxes': image_boxes[keep],
    'scores': image_scores[keep],
    'labels': image_labels[keep],
})
return detections

```

```

[ ]: def _xavier_normal_init(conv: nn.Module):
    for layer in conv.modules():
        if isinstance(layer, nn.Conv2d):
            torch.nn.init.xavier_normal_(layer.weight)
            if layer.bias is not None:
                torch.nn.init.constant_(layer.bias, 0.0)

def _kaiming_normal_init(conv: nn.Module):

```

```

for layer in conv.modules():
    if isinstance(layer, nn.Conv2d):
        torch.nn.init.kaiming_normal_(layer.weight)
        if layer.bias is not None:
            torch.nn.init.constant_(layer.bias, 0.0)

# MobileNetV3 Custom build for facemask classification
class CustomMV3(nn.Module):
    def __init__(self, base):
        super().__init__()

        self.features = base.features
        self.avgPool = base.avgpool
        self.flatten = nn.Flatten()
        self.classfier = nn.Sequential(

            nn.Linear(960, 1280, bias=True),
            nn.Hardswish(inplace=False),
            nn.Dropout(p=0.2, inplace=True),
            nn.Linear(in_features=1280, out_features=3, bias=True)

        )

    def forward(self, x):
        out = self.features(x)
        out = self.avgPool(out)
        out = self.flatten(out)
        out = self.classfier(out)
        return out

class SSDFeatureExtractorMobilenetV3(torch.nn.Module):
    """
    create 6: feature extraction layers.
    1. from base model itself from middle layers, [ Take from middle, it;s a
    ↪ experiment.]
    2. get last layer, add some conv layers to it.
    3. add 4 custom sequential layer(conv, relu, conv, relu)

    Total: 6 number of feature will be @param: returned OrderedDict("0":
    ↪ feature_map)

    """

    def __init__(self, base):

```



```

super(SSDFeatureExtractorMobilenetV3, self).__init__()

self.featuresOneFromBase = nn.Sequential(
    *base[:5]
)

self.featuresTwoFromBase = nn.Sequential(
    *base[5:10] # until InvertedResidual 9
)

self.featuresThreeFromBase = nn.Sequential(
    *base[10:] # until InvertedResidual 10 --> 16
)

"""
VGG16 backbone has put - 4 extra blocks plus 1 for high res image
MobileNetV3 backbone has put - 3 extra blocks plus 1 for 1024 res image


$$\lceil \frac{(n+2p-k)}{s} \rceil + 1$$


keeping HxW = 16x16, channels: decreasing
"""

fc1 = nn.Sequential(
    nn.Conv2d(in_channels=960, out_channels=64, kernel_size=1,
    ↪padding=0, stride=1), # FC6 with atrous
    nn.BatchNorm2d(num_features=64),
    nn.Hardswish(inplace=True),

    nn.Conv2d(in_channels=64, out_channels=128, kernel_size=1,
    ↪padding=0, stride=1), # FC6 with atrous
    nn.BatchNorm2d(num_features=128),
    nn.Hardswish(inplace=True),

    nn.Conv2d(in_channels=128, out_channels=256, kernel_size=1,
    ↪padding=0, stride=1), # FC6 with atrous
    nn.BatchNorm2d(num_features=256),

)

fc2 = nn.Sequential(
    nn.Conv2d(256, 32, kernel_size=1, padding=0, stride=2),
    nn.BatchNorm2d(num_features=32),
    nn.Hardswish(inplace=True),

    nn.Conv2d(32, 64, kernel_size=1, padding=0, stride=2),

```

```

        nn.BatchNorm2d(num_features=64),
        nn.Hardswish(inplace=True),

        nn.Conv2d(64, 128, kernel_size=1, padding=0, stride=2),
        nn.BatchNorm2d(num_features=128),
    )
    fc3 = nn.Sequential(
        nn.Conv2d(128, 64, kernel_size=1, padding=0, stride=2),
        nn.BatchNorm2d(num_features=64),
        nn.Hardswish(inplace=True),
    )
    fc4 = nn.Sequential(
        nn.Conv2d(64, 32, kernel_size=1, padding=0, stride=1),
        nn.BatchNorm2d(num_features=32),
        nn.Hardswish(inplace=True),
    )
    fc5 = nn.Sequential(
        nn.Conv2d(32, 16, kernel_size=1, padding=0, stride=2),
        nn.BatchNorm2d(num_features=16),
        nn.Hardswish(inplace=True),
    )

    _xavier_normal_init(fc1)
    _xavier_normal_init(fc2)
    _xavier_normal_init(fc3)
    _xavier_normal_init(fc4)
    _xavier_normal_init(fc5)

    self.extra = nn.ModuleList([ fc1, fc2, fc3, fc4, fc5])

    self.scale_weight = nn.Parameter(torch.ones(960) * 200)

def forward(self, x):
    output = []

    outOne = self.featuresOneFromBase(x)
    output.append(outOne)

    out = self.featuresTwoFromBase(outOne)
    output.append(out)

    out = self.featuresThreeFromBase(out)
    output.append(out)

    for i, block in enumerate(self.extra):
        out = block(out)
        output.append(out)

```



```

    ↪      [2, 3],
    ↪      [2, 3],
    ↪      [2, 3], [2]],
                                scales=[0.07, 0.15, 0.33, 0.51, 0.
    ↪69, 0.87, 1.05, 1.3, 1.5])

    print(f"number of anchors per grid for each features: {anchor_generator.
    ↪num_anchors_per_location()}")

    defaults = {
        # Rescale the input in a way compatible to the backbone

        "image_mean": [0.48235, 0.45882, 0.40784],
        "image_std": [1.0 / 255.0, 1.0 / 255.0, 1.0 / 255.0]
        # undo the 0-1 scaling of toTensor
    }

    kwargs = {**defaults}
    SSD_CUSTOM = WeightedLossSSD(backbone, anchor_generator, (size, size),
    ↪num_classes, **kwargs)

    return SSD_CUSTOM

```

```
[ ]: SSD_MODEL = get_ssd_model(num_classes=4, size=512)
```

Downloading:

"https://download.pytorch.org/models/mobilenet\_v3\_large-8738ca79.pth" to  
/root/.cache/torch/hub/checkpoints/mobilenet\_v3\_large-8738ca79.pth

0%| | 0.00/21.1M [00:00<?, ?B/s]

number of anchors per grid for each features: [6, 6, 6, 6, 6, 6, 6, 4]

```
[ ]: SSD_MODEL
```

## 11 SSD MobileNet Inbuilt

```
[ ]: SSD_MODEL = torchvision.models.detection.
    ↪ssdlite320_mobilenet_v3_large(pretrained=True)
```

## 12 Train (Inbuilt SSD)

```
[ ]: device = getDevice()
device
```

```
[ ]: device(type='cuda')
```

```
[ ]: from engine import train_one_epoch, evaluate
import utils

if torch.cuda.is_available():
    torch.cuda.empty_cache()

print(f"Training on {device}")
SSD_MODEL.to(device)

# construct an optimizer
params = [p for p in SSD_MODEL.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.01, momentum=0.9, weight_decay=0.0005)
# optimizer = torch.optim.Adam(params=params, lr=0.01, betas=(0.9, 0.009))

lr_scheduler_step_lr = torch.optim.lr_scheduler.StepLR(optimizer,
                                                         step_size=200,
                                                         gamma=0.7,)

lr_scheduler_cosine_annealing = torch.optim.lr_scheduler.
    ↪CosineAnnealingLR(optimizer, T_max=10)

# https://pytorch.org/docs/stable/generated/torch.optim.lr_scheduler.
    ↪ReduceLROnPlateau.html
lr_scheduler_platau_reduce = torch.optim.lr_scheduler.
    ↪ReduceLROnPlateau(optimizer, 'min', factor=0.01, patience=2)

num_epochs = 200

#####

loss_list = []

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    # logger = train_one_epoch(SSD_MODEL, optimizer, train_data_loader, device,
    ↪epoch, print_freq=100)
    # TRAINING BATCH OF DATA
    epoch_loss = []
    SSD_MODEL.train()
```

```

for i, data in enumerate(train_data_loader):

    # data of each batch
    images, targets = data

    # copy data to cuda
    inputs = list(image.to(device) for image in images)
    targets = [{k: v.to(device) for k, v in t.items()} for t in targets]

    # set gradients to 0
    optimizer.zero_grad()

    # pass input to model
    loss_dict = SSD_MODEL(inputs, targets)

    # loss
    losses = sum(loss for loss in loss_dict.values())
    epoch_loss.append(losses.item())

    # backprop
    losses.backward()

    # update weights
    optimizer.step()
    #-----#

    # update the learning rate
    # lr_scheduler_step_lr.step()

    # Epoch end
    train_loss_epoch = np.mean(epoch_loss)
    print("\n")
    print(f"epoch: {epoch}, training loss: {train_loss_epoch}, lr: {optimizer.
→param_groups[0]['lr']} ") # shows every training loss and parameters
    loss_list.append(train_loss_epoch)

    if epoch > 100:
        if loss_list[epoch] > loss_list[epoch-10]:

            torch.save({
                'epoch': epoch,
                'model_state_dict': SSD_MODEL.state_dict(),
                'optimizer_state_dict': optimizer.state_dict(),
            }, "/content/gdrive/MyDrive/facemask_detection/save_model/
→facemaskDetectionSSD_320x320_38MAP.pth")

```

```

        break

    # evaluate on the test dataset
    coco_evaluator = evaluate(SSD_MODEL, valid_data_loader, device=device)

    # reduce lr on platau, after validation
    # lr_scheduler_platau_reduce.step(train_loss_epoch)

    print(coco_evaluator.coco_eval)

    torch.save({
        'epoch': epoch,
        'model_state_dict': SSD_MODEL.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
    }, "/content/gdrive/MyDrive/facemask_detection/save_model/
→facemaskDetectionSSD_320x320_38MAP.pth")

    if epoch > 30:
        optimizer.param_groups[0]['lr'] = 0.001 * 0.1

```

Training on cuda

```

Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at
0x7f698179b050>
Traceback (most recent call last):
  File "/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py",
line 1328, in __del__
    self._shutdown_workers()
  File "/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py",
line 1320, in _shutdown_workers
    if w.is_alive():
  File "/usr/lib/python3.7/multiprocessing/process.py", line 151, in is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'
AssertionError: can only test a child process
Exception ignored in: <function _MultiProcessingDataLoaderIter.__del__ at
0x7f698179b050>
Traceback (most recent call last):
  File "/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py",
line 1328, in __del__
    self._shutdown_workers()
  File "/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py",
line 1320, in _shutdown_workers
    if w.is_alive():

```

```
File "/usr/lib/python3.7/multiprocessing/process.py", line 151, in is_alive
    assert self._parent_pid == os.getpid(), 'can only test a child process'
AssertionError: can only test a child process
```

```
epoch: 0, training loss: 3.186100948601961, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:49 model_time: 1.1894 (1.1894) evaluator_time:
0.2984 (0.2984) time: 2.2656 data: 0.7615 max mem: 5550
Test: [21/22] eta: 0:00:01 model_time: 0.7167 (0.7430) evaluator_time:
0.1709 (0.2303) time: 0.9663 data: 0.0672 max mem: 5550
Test: Total time: 0:00:24 (1.1024 s / it)
Averaged stats: model_time: 0.7167 (0.7430) evaluator_time: 0.1709 (0.2303)
Accumulating evaluation results...
DONE (t=0.76s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.246
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.433
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.253
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.025
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.222
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.453
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.270
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.410
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.452
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.137
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.457
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.692
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69735b8650>}
```

```
epoch: 1, training loss: 2.359612367115915, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:01:04 model_time: 1.2637 (1.2637) evaluator_time:
0.3787 (0.3787) time: 2.9462 data: 1.2812 max mem: 5550
Test: [21/22] eta: 0:00:01 model_time: 0.6775 (0.7108) evaluator_time:
0.1803 (0.2358) time: 0.9811 data: 0.0604 max mem: 5550
Test: Total time: 0:00:23 (1.0833 s / it)
Averaged stats: model_time: 0.6775 (0.7108) evaluator_time: 0.1803 (0.2358)
Accumulating evaluation results...
DONE (t=0.81s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.293
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.493
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.321
```



```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.053
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.265
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.521
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.276
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.447
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.482
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.181
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.492
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.700
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6962ade050>}

```

epoch: 2, training loss: 2.0934839230030775, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:01:11 model\_time: 1.4479 (1.4479) evaluator\_time: 0.5906 (0.5906) time: 3.2723 data: 1.2170 max mem: 5550

Test: [21/22] eta: 0:00:01 model\_time: 0.6975 (0.7349) evaluator\_time: 0.1850 (0.2524) time: 0.9968 data: 0.0588 max mem: 5550

Test: Total time: 0:00:24 (1.1180 s / it)

Averaged stats: model\_time: 0.6975 (0.7349) evaluator\_time: 0.1850 (0.2524)

Accumulating evaluation results...

DONE (t=0.92s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.292
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.517
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.303
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.044
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.282
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.495
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.310
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.465
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.501
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.204
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.504
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.732
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6972344350>}

```

epoch: 3, training loss: 1.9617709023877978, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:43 model\_time: 0.9125 (0.9125) evaluator\_time: 0.2731 (0.2731) time: 1.9692 data: 0.7736 max mem: 5550

Test: [21/22] eta: 0:00:01 model\_time: 0.6891 (0.7058) evaluator\_time: 0.2238 (0.2531) time: 1.0332 data: 0.0652 max mem: 5550

Test: Total time: 0:00:23 (1.0731 s / it)

Averaged stats: model\_time: 0.6891 (0.7058) evaluator\_time: 0.2238 (0.2531)

Accumulating evaluation results...

DONE (t=0.83s).

IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.302
Average Precision	(AP)	@[ IoU=0.50	area=	all	maxDets=100 ]	=	0.526
Average Precision	(AP)	@[ IoU=0.75	area=	all	maxDets=100 ]	=	0.300
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.047
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.278
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.537
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 1 ]	=	0.293
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 10 ]	=	0.455
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.490
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.221
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.497
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.689

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697234fe10>}

epoch: 4, training loss: 1.8344742273911834, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:56 model\_time: 0.9836 (0.9836) evaluator\_time:  
0.2860 (0.2860) time: 2.5598 data: 1.2804 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5733 (0.5983) evaluator\_time:  
0.1836 (0.2425) time: 0.8733 data: 0.0603 max mem: 5550

Test: Total time: 0:00:21 (0.9711 s / it)

Averaged stats: model\_time: 0.5733 (0.5983) evaluator\_time: 0.1836 (0.2425)

Accumulating evaluation results...

DONE (t=0.69s).

IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.310
Average Precision	(AP)	@[ IoU=0.50	area=	all	maxDets=100 ]	=	0.531
Average Precision	(AP)	@[ IoU=0.75	area=	all	maxDets=100 ]	=	0.299
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.062
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.306
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.501
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 1 ]	=	0.294
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 10 ]	=	0.435
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.460
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.222
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.461
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.647

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6972395410>}

epoch: 5, training loss: 1.71666428912431, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:01:07 model\_time: 1.0615 (1.0615) evaluator\_time: 0.6913 (0.6913) time: 3.0752 data: 1.3053 max mem: 5550  
Test: [21/22] eta: 0:00:01 model\_time: 0.6565 (0.6466) evaluator\_time: 0.2036 (0.2584) time: 0.9278 data: 0.0569 max mem: 5550  
Test: Total time: 0:00:22 (1.0420 s / it)  
Averaged stats: model\_time: 0.6565 (0.6466) evaluator\_time: 0.2036 (0.2584)  
Accumulating evaluation results...  
DONE (t=0.82s).  
IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area= all   maxDets=100 ]	= 0.328
Average Precision	(AP)	@[ IoU=0.50	area= all   maxDets=100 ]	= 0.558
Average Precision	(AP)	@[ IoU=0.75	area= all   maxDets=100 ]	= 0.340
Average Precision	(AP)	@[ IoU=0.50:0.95	area= small   maxDets=100 ]	= 0.056
Average Precision	(AP)	@[ IoU=0.50:0.95	area= medium   maxDets=100 ]	= 0.324
Average Precision	(AP)	@[ IoU=0.50:0.95	area= large   maxDets=100 ]	= 0.535
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all   maxDets= 1 ]	= 0.313
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all   maxDets= 10 ]	= 0.482
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all   maxDets=100 ]	= 0.519
Average Recall	(AR)	@[ IoU=0.50:0.95	area= small   maxDets=100 ]	= 0.264
Average Recall	(AR)	@[ IoU=0.50:0.95	area= medium   maxDets=100 ]	= 0.542
Average Recall	(AR)	@[ IoU=0.50:0.95	area= large   maxDets=100 ]	= 0.680

{'bbox': <pycocotools.cocoeval.COCDeval object at 0x7f6972395590>}

epoch: 6, training loss: 1.6804218059405684, lr: 0.01  
creating index...  
index created!

Test: [ 0/22] eta: 0:00:50 model\_time: 0.9336 (0.9336) evaluator\_time: 0.4453 (0.4453) time: 2.2803 data: 0.8815 max mem: 5550  
Test: [21/22] eta: 0:00:01 model\_time: 0.6100 (0.6563) evaluator\_time: 0.2080 (0.2549) time: 0.9292 data: 0.0608 max mem: 5550  
Test: Total time: 0:00:22 (1.0271 s / it)  
Averaged stats: model\_time: 0.6100 (0.6563) evaluator\_time: 0.2080 (0.2549)  
Accumulating evaluation results...  
DONE (t=0.80s).  
IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area= all   maxDets=100 ]	= 0.339
Average Precision	(AP)	@[ IoU=0.50	area= all   maxDets=100 ]	= 0.576
Average Precision	(AP)	@[ IoU=0.75	area= all   maxDets=100 ]	= 0.355
Average Precision	(AP)	@[ IoU=0.50:0.95	area= small   maxDets=100 ]	= 0.096
Average Precision	(AP)	@[ IoU=0.50:0.95	area= medium   maxDets=100 ]	= 0.328
Average Precision	(AP)	@[ IoU=0.50:0.95	area= large   maxDets=100 ]	= 0.549
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all   maxDets= 1 ]	= 0.334
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all   maxDets= 10 ]	= 0.489
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all   maxDets=100 ]	= 0.527
Average Recall	(AR)	@[ IoU=0.50:0.95	area= small   maxDets=100 ]	= 0.239
Average Recall	(AR)	@[ IoU=0.50:0.95	area= medium   maxDets=100 ]	= 0.530
Average Recall	(AR)	@[ IoU=0.50:0.95	area= large   maxDets=100 ]	= 0.748

```
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f697234f5d0>}
```

epoch: 7, training loss: 1.6146570211276412, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:01:21 model\_time: 0.9658 (0.9658) evaluator\_time:  
0.6012 (0.6012) time: 3.6845 data: 2.0914 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5532 (0.5952) evaluator\_time:  
0.2232 (0.2335) time: 0.8651 data: 0.0610 max mem: 5550

Test: Total time: 0:00:21 (0.9977 s / it)

Averaged stats: model\_time: 0.5532 (0.5952) evaluator\_time: 0.2232 (0.2335)

Accumulating evaluation results...

DONE (t=0.74s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.313

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.530

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.334

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.072

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.284

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.539

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.299

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.466

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.498

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.225

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.497

Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.719

```
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6962ade0d0>}
```

epoch: 8, training loss: 1.5138298412784934, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:01:01 model\_time: 0.5983 (0.5983) evaluator\_time:  
0.3235 (0.3235) time: 2.8097 data: 1.8782 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5724 (0.5750) evaluator\_time:  
0.2098 (0.2449) time: 0.8749 data: 0.0617 max mem: 5550

Test: Total time: 0:00:21 (0.9779 s / it)

Averaged stats: model\_time: 0.5724 (0.5750) evaluator\_time: 0.2098 (0.2449)

Accumulating evaluation results...

DONE (t=0.77s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.330

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.565

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.358

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.048

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.296

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.592

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.308
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.457
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.502
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.175
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.535
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.700
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6973569ad0>}

```

```

epoch: 9, training loss: 1.490667070262134, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:01:04 model_time: 0.5753 (0.5753) evaluator_time:
0.3841 (0.3841) time: 2.9528 data: 1.9804 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.5953 (0.5715) evaluator_time:
0.2564 (0.2591) time: 0.8890 data: 0.0608 max mem: 5550
Test: Total time: 0:00:21 (0.9929 s / it)
Averaged stats: model_time: 0.5953 (0.5715) evaluator_time: 0.2564 (0.2591)
Accumulating evaluation results...
DONE (t=0.80s).
IoU metric: bbox

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.338
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.595
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.337
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.062
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.318
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.591
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.306
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.453
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.482
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.231
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.455
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.731
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6962784dd0>}

```

```

epoch: 10, training loss: 1.4149223840795457, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:53 model_time: 1.0953 (1.0953) evaluator_time:
0.5397 (0.5397) time: 2.4366 data: 0.7841 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.5084 (0.5634) evaluator_time:
0.2272 (0.2566) time: 0.8336 data: 0.0664 max mem: 5550
Test: Total time: 0:00:20 (0.9397 s / it)
Averaged stats: model_time: 0.5084 (0.5634) evaluator_time: 0.2272 (0.2566)
Accumulating evaluation results...
DONE (t=0.74s).
IoU metric: bbox

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.340
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.567
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.366
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.057
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.342
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.556
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.320
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.481
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.509
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.238
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.492
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.754
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f697219fe90>}

```

```

epoch: 11, training loss: 1.3522069291211665, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:53 model_time: 1.0701 (1.0701) evaluator_time:
0.5136 (0.5136) time: 2.4098 data: 0.8082 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.5842 (0.6257) evaluator_time:
0.1885 (0.2493) time: 0.8945 data: 0.0642 max mem: 5550
Test: Total time: 0:00:21 (0.9915 s / it)
Averaged stats: model_time: 0.5842 (0.6257) evaluator_time: 0.1885 (0.2493)
Accumulating evaluation results...
DONE (t=0.81s).

```

```

IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.347
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.579
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.363
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.072
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.331
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.580
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.333
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.491
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.538
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.239
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.561
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.731
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69711c2210>}

```

```

epoch: 12, training loss: 1.3577878116630018, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:46 model_time: 0.9629 (0.9629) evaluator_time:
0.2134 (0.2134) time: 2.1229 data: 0.9285 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.5563 (0.6085) evaluator_time:

```

```

0.2072 (0.2530)  time: 0.9142  data: 0.0650  max mem: 5550
Test: Total time: 0:00:21 (0.9825 s / it)
Averaged stats: model_time: 0.5563 (0.6085)  evaluator_time: 0.2072 (0.2530)
Accumulating evaluation results...
DONE (t=0.79s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.349
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.573
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.382
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.061
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.344
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.574
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.319
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.489
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.540
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.254
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.575
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.701
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69626d9110>}

```

```

epoch: 13, training loss: 1.309306944720447, lr: 0.01
creating index...
index created!
Test: [ 0/22]  eta: 0:00:43  model_time: 0.7828 (0.7828)  evaluator_time:
0.4282 (0.4282)  time: 1.9839  data: 0.7570  max mem: 5550
Test: [21/22]  eta: 0:00:00  model_time: 0.4709 (0.5370)  evaluator_time:
0.2094 (0.2459)  time: 0.8389  data: 0.0681  max mem: 5550
Test: Total time: 0:00:19 (0.9010 s / it)
Averaged stats: model_time: 0.4709 (0.5370)  evaluator_time: 0.2094 (0.2459)
Accumulating evaluation results...
DONE (t=0.65s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.297
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.509
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.309
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.062
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.277
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.510
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.311
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.463
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.492
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.206
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.492
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.720
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69661bc750>}

```

```

epoch: 14, training loss: 1.2380990898236632, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:50 model_time: 0.4606 (0.4606) evaluator_time:
0.3363 (0.3363) time: 2.2732 data: 1.4666 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.5109 (0.5349) evaluator_time:
0.2191 (0.2518) time: 0.8496 data: 0.0627 max mem: 5550
Test: Total time: 0:00:20 (0.9286 s / it)
Averaged stats: model_time: 0.5109 (0.5349) evaluator_time: 0.2191 (0.2518)
Accumulating evaluation results...
DONE (t=0.73s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.315
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.543
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.323
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.055
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.301
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.541
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.309
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.460
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.495
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.229
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.478
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.736
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f696f1d5bd0>}

```

```

epoch: 15, training loss: 1.2387789157219231, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:47 model_time: 0.8071 (0.8071) evaluator_time:
0.4555 (0.4555) time: 2.1471 data: 0.8741 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4823 (0.5552) evaluator_time:
0.1883 (0.2603) time: 0.8431 data: 0.0641 max mem: 5550
Test: Total time: 0:00:20 (0.9317 s / it)
Averaged stats: model_time: 0.4823 (0.5552) evaluator_time: 0.1883 (0.2603)
Accumulating evaluation results...
DONE (t=0.76s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.364
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.601
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.413
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.074
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.349
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.612
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.330
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.505
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.527

```



```
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.257
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.517
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.758
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69628efe10>}
```

epoch: 16, training loss: 1.2595119569450617, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:50 model\_time: 1.0183 (1.0183) evaluator\_time:  
0.3202 (0.3202) time: 2.3113 data: 0.9556 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5156 (0.5521) evaluator\_time:  
0.3000 (0.2827) time: 0.8773 data: 0.0689 max mem: 5550

Test: Total time: 0:00:21 (0.9586 s / it)

Averaged stats: model\_time: 0.5156 (0.5521) evaluator\_time: 0.3000 (0.2827)

Accumulating evaluation results...

DONE (t=0.79s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.345
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.585
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.072
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.350
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.542
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.327
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.502
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.532
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.266
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.521
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.759
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697234f090>}
```

epoch: 17, training loss: 1.170159401372075, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:01:05 model\_time: 0.5757 (0.5757) evaluator\_time:  
0.2992 (0.2992) time: 2.9855 data: 2.1010 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5201 (0.5638) evaluator\_time:  
0.2068 (0.2241) time: 0.8311 data: 0.0587 max mem: 5550

Test: Total time: 0:00:20 (0.9539 s / it)

Averaged stats: model\_time: 0.5201 (0.5638) evaluator\_time: 0.2068 (0.2241)

Accumulating evaluation results...

DONE (t=0.69s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.357
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.602
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.375
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.072
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.353
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.588
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.319
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.504
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.523
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.233
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.528
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.742
{'bbox': <pycocotools.cocoeval.COC0eval object at 0x7f6961fef2d0>}

```

epoch: 18, training loss: 1.1530232704244554, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:51 model\_time: 0.8195 (0.8195) evaluator\_time: 0.5892 (0.5892) time: 2.3456 data: 0.9192 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5223 (0.5672) evaluator\_time: 0.2176 (0.2682) time: 0.8886 data: 0.0665 max mem: 5550

Test: Total time: 0:00:21 (0.9608 s / it)

Averaged stats: model\_time: 0.5223 (0.5672) evaluator\_time: 0.2176 (0.2682)

Accumulating evaluation results...

DONE (t=0.73s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.371
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.613
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.403
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.071
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.348
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.621
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.318
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.497
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.526
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.267
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.516
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.747
{'bbox': <pycocotools.cocoeval.COC0eval object at 0x7f69628e0e10>}

```

epoch: 19, training loss: 1.123286777175963, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:01:09 model\_time: 0.5776 (0.5776) evaluator\_time: 0.3394 (0.3394) time: 3.1705 data: 2.2375 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4908 (0.4963) evaluator\_time: 0.2199 (0.2508) time: 0.7928 data: 0.0625 max mem: 5550

Test: Total time: 0:00:20 (0.9232 s / it)

Averaged stats: model\_time: 0.4908 (0.4963) evaluator\_time: 0.2199 (0.2508)

Accumulating evaluation results...

DONE (t=0.72s).

IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.321
Average Precision	(AP)	@[ IoU=0.50	area=	all	maxDets=100 ]	=	0.561
Average Precision	(AP)	@[ IoU=0.75	area=	all	maxDets=100 ]	=	0.315
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.086
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.314
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.553
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 1 ]	=	0.320
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 10 ]	=	0.489
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.519
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.334
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.498
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.702

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6971255550>}

epoch: 20, training loss: 1.0867871632799506, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:01:32 model\_time: 0.5654 (0.5654) evaluator\_time: 0.4584 (0.4584) time: 4.1992 data: 3.1653 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5318 (0.5200) evaluator\_time: 0.1789 (0.2166) time: 0.7907 data: 0.0560 max mem: 5550

Test: Total time: 0:00:20 (0.9469 s / it)

Averaged stats: model\_time: 0.5318 (0.5200) evaluator\_time: 0.1789 (0.2166)

Accumulating evaluation results...

DONE (t=0.69s).

IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.349
Average Precision	(AP)	@[ IoU=0.50	area=	all	maxDets=100 ]	=	0.583
Average Precision	(AP)	@[ IoU=0.75	area=	all	maxDets=100 ]	=	0.382
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.090
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.329
Average Precision	(AP)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.597
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 1 ]	=	0.329
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets= 10 ]	=	0.473
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	all	maxDets=100 ]	=	0.521
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	small	maxDets=100 ]	=	0.288
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	medium	maxDets=100 ]	=	0.509
Average Recall	(AR)	@[ IoU=0.50:0.95	area=	large	maxDets=100 ]	=	0.727

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6971257610>}

epoch: 21, training loss: 1.0863277269527316, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:47 model\_time: 0.8751 (0.8751) evaluator\_time: 0.3930 (0.3930) time: 2.1780 data: 0.8905 max mem: 5550  
 Test: [21/22] eta: 0:00:00 model\_time: 0.4928 (0.5344) evaluator\_time: 0.1895 (0.2384) time: 0.8052 data: 0.0619 max mem: 5550  
 Test: Total time: 0:00:19 (0.8891 s / it)  
 Averaged stats: model\_time: 0.4928 (0.5344) evaluator\_time: 0.1895 (0.2384)  
 Accumulating evaluation results...  
 DONE (t=0.63s).  
 IoU metric: bbox

Average Precision	(AP) @[ IoU=0.50:0.95   area=	all   maxDets=100 ]	= 0.346
Average Precision	(AP) @[ IoU=0.50   area=	all   maxDets=100 ]	= 0.579
Average Precision	(AP) @[ IoU=0.75   area=	all   maxDets=100 ]	= 0.361
Average Precision	(AP) @[ IoU=0.50:0.95   area=	small   maxDets=100 ]	= 0.073
Average Precision	(AP) @[ IoU=0.50:0.95   area=	medium   maxDets=100 ]	= 0.327
Average Precision	(AP) @[ IoU=0.50:0.95   area=	large   maxDets=100 ]	= 0.583
Average Recall	(AR) @[ IoU=0.50:0.95   area=	all   maxDets= 1 ]	= 0.309
Average Recall	(AR) @[ IoU=0.50:0.95   area=	all   maxDets= 10 ]	= 0.469
Average Recall	(AR) @[ IoU=0.50:0.95   area=	all   maxDets=100 ]	= 0.492
Average Recall	(AR) @[ IoU=0.50:0.95   area=	small   maxDets=100 ]	= 0.241
Average Recall	(AR) @[ IoU=0.50:0.95   area=	medium   maxDets=100 ]	= 0.474
Average Recall	(AR) @[ IoU=0.50:0.95   area=	large   maxDets=100 ]	= 0.725

{'bbox': <pycocotools.cocoeval.COCDeval object at 0x7f6a8ef6d750>}

epoch: 22, training loss: 1.0624138824641705, lr: 0.01  
 creating index...  
 index created!

Test: [ 0/22] eta: 0:01:09 model\_time: 0.6297 (0.6297) evaluator\_time: 0.5183 (0.5183) time: 3.1746 data: 2.0096 max mem: 5550  
 Test: [21/22] eta: 0:00:00 model\_time: 0.4672 (0.5393) evaluator\_time: 0.1980 (0.2375) time: 0.8187 data: 0.0621 max mem: 5550  
 Test: Total time: 0:00:20 (0.9420 s / it)  
 Averaged stats: model\_time: 0.4672 (0.5393) evaluator\_time: 0.1980 (0.2375)  
 Accumulating evaluation results...  
 DONE (t=0.71s).  
 IoU metric: bbox

Average Precision	(AP) @[ IoU=0.50:0.95   area=	all   maxDets=100 ]	= 0.366
Average Precision	(AP) @[ IoU=0.50   area=	all   maxDets=100 ]	= 0.596
Average Precision	(AP) @[ IoU=0.75   area=	all   maxDets=100 ]	= 0.415
Average Precision	(AP) @[ IoU=0.50:0.95   area=	small   maxDets=100 ]	= 0.080
Average Precision	(AP) @[ IoU=0.50:0.95   area=	medium   maxDets=100 ]	= 0.364
Average Precision	(AP) @[ IoU=0.50:0.95   area=	large   maxDets=100 ]	= 0.569
Average Recall	(AR) @[ IoU=0.50:0.95   area=	all   maxDets= 1 ]	= 0.323
Average Recall	(AR) @[ IoU=0.50:0.95   area=	all   maxDets= 10 ]	= 0.492
Average Recall	(AR) @[ IoU=0.50:0.95   area=	all   maxDets=100 ]	= 0.527
Average Recall	(AR) @[ IoU=0.50:0.95   area=	small   maxDets=100 ]	= 0.291
Average Recall	(AR) @[ IoU=0.50:0.95   area=	medium   maxDets=100 ]	= 0.520
Average Recall	(AR) @[ IoU=0.50:0.95   area=	large   maxDets=100 ]	= 0.723

{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6962784b90>}

epoch: 23, training loss: 1.0344882323406637, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:44 model\_time: 0.7480 (0.7480) evaluator\_time:  
0.3626 (0.3626) time: 2.0202 data: 0.8834 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5180 (0.5205) evaluator\_time:  
0.2069 (0.2315) time: 0.7856 data: 0.0563 max mem: 5550

Test: Total time: 0:00:20 (0.9212 s / it)

Averaged stats: model\_time: 0.5180 (0.5205) evaluator\_time: 0.2069 (0.2315)

Accumulating evaluation results...

DONE (t=0.67s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.347

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.578

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.349

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.074

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.355

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.550

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.322

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.498

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.541

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.294

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.527

Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.757

{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f696209d390>}

epoch: 24, training loss: 1.073603137396276, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:51 model\_time: 0.7724 (0.7724) evaluator\_time:  
0.2871 (0.2871) time: 2.3367 data: 1.2612 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4641 (0.5188) evaluator\_time:  
0.2173 (0.2392) time: 0.8011 data: 0.0632 max mem: 5550

Test: Total time: 0:00:19 (0.8919 s / it)

Averaged stats: model\_time: 0.4641 (0.5188) evaluator\_time: 0.2173 (0.2392)

Accumulating evaluation results...

DONE (t=0.64s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.335

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.560

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.365

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.318

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.579

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.343
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.495
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.531
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.273
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.526
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.743
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6962bbe8d0>}

```

epoch: 25, training loss: 1.0194266475737095, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:51 model\_time: 0.8765 (0.8765) evaluator\_time:  
0.5664 (0.5664) time: 2.3333 data: 0.8724 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4668 (0.5326) evaluator\_time:  
0.2082 (0.2652) time: 0.8114 data: 0.0583 max mem: 5550

Test: Total time: 0:00:20 (0.9223 s / it)

Averaged stats: model\_time: 0.4668 (0.5326) evaluator\_time: 0.2082 (0.2652)

Accumulating evaluation results...

DONE (t=0.66s).

IoU metric: bbox

```

Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.348
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.577
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.367
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.079
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.346
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.567
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.320
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.462
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.530
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.328
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.542
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.667

```

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69710c02d0>}

epoch: 26, training loss: 0.9851195416413248, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:47 model\_time: 0.7315 (0.7315) evaluator\_time:  
0.4646 (0.4646) time: 2.1713 data: 0.9579 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4635 (0.5369) evaluator\_time:  
0.2316 (0.2556) time: 0.8610 data: 0.0681 max mem: 5550

Test: Total time: 0:00:20 (0.9176 s / it)

Averaged stats: model\_time: 0.4635 (0.5369) evaluator\_time: 0.2316 (0.2556)

Accumulating evaluation results...

DONE (t=0.75s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.339
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.561
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.086
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.331
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.575
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.324
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.478
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.542
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.319
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.528
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739
{'bbox': <pycocotools.cocoeval.COCDeval object at 0x7f69713cdc50>}

```

epoch: 27, training loss: 0.9502438558265567, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:45 model\_time: 0.8496 (0.8496) evaluator\_time:  
0.4357 (0.4357) time: 2.0837 data: 0.7772 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5175 (0.5465) evaluator\_time:  
0.1929 (0.2482) time: 0.8403 data: 0.0668 max mem: 5550

Test: Total time: 0:00:20 (0.9174 s / it)

Averaged stats: model\_time: 0.5175 (0.5465) evaluator\_time: 0.1929 (0.2482)

Accumulating evaluation results...

DONE (t=0.63s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.375
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.614
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.427
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.078
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.336
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.645
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.350
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.511
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.559
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.288
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.563
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.764
{'bbox': <pycocotools.cocoeval.COCDeval object at 0x7f6962bbefd0>}

```

epoch: 28, training loss: 0.9401585757732391, lr: 0.01

creating index...

index created!

Test: [ 0/22] eta: 0:00:57 model\_time: 0.4816 (0.4816) evaluator\_time:  
0.2620 (0.2620) time: 2.6205 data: 1.8670 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4707 (0.4990) evaluator\_time:

```

0.2238 (0.2356)  time: 0.7922  data: 0.0624  max mem: 5550
Test: Total time: 0:00:19 (0.8939 s / it)
Averaged stats: model_time: 0.4707 (0.4990)  evaluator_time: 0.2238 (0.2356)
Accumulating evaluation results...
DONE (t=0.58s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.360
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.577
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.379
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.079
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.360
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.551
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.350
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.498
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.533
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.286
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.538
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.715
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697112b3d0>}

```

```

epoch: 29, training loss: 0.992657536175102, lr: 0.01
creating index...
index created!
Test: [ 0/22]  eta: 0:00:41  model_time: 0.8072 (0.8072)  evaluator_time:
0.2491 (0.2491)  time: 1.8664  data: 0.7921  max mem: 5550
Test: [21/22]  eta: 0:00:00  model_time: 0.5005 (0.5326)  evaluator_time:
0.2079 (0.2344)  time: 0.8206  data: 0.0781  max mem: 5550
Test: Total time: 0:00:19 (0.8960 s / it)
Averaged stats: model_time: 0.5005 (0.5326)  evaluator_time: 0.2079 (0.2344)
Accumulating evaluation results...
DONE (t=0.59s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.334
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.546
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.355
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.065
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.311
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.575
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.324
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.461
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.503
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.277
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.507
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.672
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697109b1d0>}

```



```

epoch: 30, training loss: 0.9635079335421324, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:48 model_time: 0.8621 (0.8621) evaluator_time:
0.2282 (0.2282) time: 2.1888 data: 1.0762 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4928 (0.5240) evaluator_time:
0.2602 (0.2552) time: 0.8309 data: 0.0658 max mem: 5550
Test: Total time: 0:00:20 (0.9091 s / it)
Averaged stats: model_time: 0.4928 (0.5240) evaluator_time: 0.2602 (0.2552)
Accumulating evaluation results...
DONE (t=0.68s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.343
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.554
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.357
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.086
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.345
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.536
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.337
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.506
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.554
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.334
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.697
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69735acf90>}

```

```

epoch: 31, training loss: 0.9568158886395395, lr: 0.01
creating index...
index created!
Test: [ 0/22] eta: 0:00:52 model_time: 0.7099 (0.7099) evaluator_time:
0.4941 (0.4941) time: 2.4026 data: 1.1792 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4442 (0.5151) evaluator_time:
0.1962 (0.2197) time: 0.7633 data: 0.0682 max mem: 5550
Test: Total time: 0:00:19 (0.8704 s / it)
Averaged stats: model_time: 0.4442 (0.5151) evaluator_time: 0.1962 (0.2197)
Accumulating evaluation results...
DONE (t=0.60s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.324
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.532
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.364
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.059
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.281
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.605
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.314
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.463
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.516

```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.242
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.529
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.712
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6967058310>}
```

epoch: 32, training loss: 0.8973044506274164, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:17 model\_time: 0.6748 (0.6748) evaluator\_time:  
0.4370 (0.4370) time: 3.5086 data: 2.3870 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4804 (0.5052) evaluator\_time:  
0.1564 (0.2183) time: 0.7810 data: 0.0658 max mem: 5550

Test: Total time: 0:00:20 (0.9108 s / it)

Averaged stats: model\_time: 0.4804 (0.5052) evaluator\_time: 0.1564 (0.2183)

Accumulating evaluation results...

DONE (t=0.61s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.370
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.604
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.421
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.082
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.355
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.618
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.355
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.520
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.548
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.322
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.542
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.734
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6973560fd0>}
```

epoch: 33, training loss: 0.8478483376093209, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:54 model\_time: 0.9338 (0.9338) evaluator\_time:  
0.6664 (0.6664) time: 2.4605 data: 0.8425 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4896 (0.5176) evaluator\_time:  
0.2115 (0.2608) time: 0.8210 data: 0.0686 max mem: 5550

Test: Total time: 0:00:19 (0.9035 s / it)

Averaged stats: model\_time: 0.4896 (0.5176) evaluator\_time: 0.2115 (0.2608)

Accumulating evaluation results...

DONE (t=0.62s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.379
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.616
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.430
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.085
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.365
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.620
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.356
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.523
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.550
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.319
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.547
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69722ed050>}

```

epoch: 34, training loss: 0.8454588232561946, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:24 model\_time: 0.6664 (0.6664) evaluator\_time:  
0.5725 (0.5725) time: 3.8198 data: 2.5632 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4774 (0.4927) evaluator\_time:  
0.1871 (0.2307) time: 0.7695 data: 0.0696 max mem: 5550

Test: Total time: 0:00:20 (0.9201 s / it)

Averaged stats: model\_time: 0.4774 (0.4927) evaluator\_time: 0.1871 (0.2307)

Accumulating evaluation results...

DONE (t=0.64s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.380
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.619
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.430
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.088
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.368
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.617
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.357
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.523
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.550
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.317
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.548
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.736
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69712028d0>}

```

epoch: 35, training loss: 0.8241126104258001, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:52 model\_time: 0.8398 (0.8398) evaluator\_time:  
0.7029 (0.7029) time: 2.4029 data: 0.8405 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4699 (0.5199) evaluator\_time:  
0.1719 (0.2464) time: 0.7927 data: 0.0674 max mem: 5550

Test: Total time: 0:00:19 (0.8886 s / it)

Averaged stats: model\_time: 0.4699 (0.5199) evaluator\_time: 0.1719 (0.2464)

Accumulating evaluation results...

DONE (t=0.62s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.383
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.621
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.436
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.091
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.370
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.617
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.360
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.526
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.552
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.313
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.551
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.741
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6971207810>}
```

epoch: 36, training loss: 0.8183679520152509, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:50 model\_time: 0.7588 (0.7588) evaluator\_time:  
0.4941 (0.4941) time: 2.2865 data: 1.0235 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4681 (0.5095) evaluator\_time:  
0.1829 (0.2376) time: 0.7852 data: 0.0621 max mem: 5550

Test: Total time: 0:00:20 (0.9145 s / it)

Averaged stats: model\_time: 0.4681 (0.5095) evaluator\_time: 0.1829 (0.2376)

Accumulating evaluation results...

DONE (t=0.64s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.625
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.436
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.092
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.371
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.619
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.358
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.526
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.564
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.313
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.576
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6972041290>}
```

epoch: 37, training loss: 0.78047746559605, lr: 0.0001

creating index...

index created!

```

Test: [ 0/22] eta: 0:01:04 model_time: 0.5404 (0.5404) evaluator_time:
0.4253 (0.4253) time: 2.9338 data: 1.9575 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4823 (0.5013) evaluator_time:
0.2182 (0.2481) time: 0.8109 data: 0.0606 max mem: 5550
Test: Total time: 0:00:20 (0.9140 s / it)
Averaged stats: model_time: 0.4823 (0.5013) evaluator_time: 0.2182 (0.2481)
Accumulating evaluation results...
DONE (t=0.60s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.385
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.627
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.437
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.092
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.371
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.618
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.350
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.520
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.559
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.314
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.567
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.736
{'bbox': <pycocotools.cocoeval.COCDeval object at 0x7f69735824d0>}

```

```

epoch: 38, training loss: 0.775688661262393, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:43 model_time: 0.7959 (0.7959) evaluator_time:
0.3981 (0.3981) time: 1.9843 data: 0.7730 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4941 (0.5247) evaluator_time:
0.2286 (0.2587) time: 0.8050 data: 0.0673 max mem: 5550
Test: Total time: 0:00:19 (0.9036 s / it)
Averaged stats: model_time: 0.4941 (0.5247) evaluator_time: 0.2286 (0.2587)
Accumulating evaluation results...
DONE (t=0.61s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.386
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.435
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.092
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.374
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.617
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.352
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.517
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.555
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.316
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.560
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.736

```

{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69672a5f50>}

epoch: 39, training loss: 0.7727686730213463, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:55 model\_time: 1.0758 (1.0758) evaluator\_time:  
0.4522 (0.4522) time: 2.5221 data: 0.9771 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4517 (0.5428) evaluator\_time:  
0.1946 (0.2573) time: 0.7979 data: 0.0598 max mem: 5550

Test: Total time: 0:00:20 (0.9175 s / it)

Averaged stats: model\_time: 0.4517 (0.5428) evaluator\_time: 0.1946 (0.2573)

Accumulating evaluation results...

DONE (t=0.67s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.385

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.629

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.430

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.375

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.618

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.348

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.510

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.557

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.343

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.548

Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.741

{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f697359ea50>}

epoch: 40, training loss: 0.7794302888214588, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:37 model\_time: 0.6913 (0.6913) evaluator\_time:  
0.1953 (0.1953) time: 1.7182 data: 0.8138 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4646 (0.5191) evaluator\_time:  
0.2434 (0.2426) time: 0.8056 data: 0.0656 max mem: 5550

Test: Total time: 0:00:19 (0.8798 s / it)

Averaged stats: model\_time: 0.4646 (0.5191) evaluator\_time: 0.2434 (0.2426)

Accumulating evaluation results...

DONE (t=0.61s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.386

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.628

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.433

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.371

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.620

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.350
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.507
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.551
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.324
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.540
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.749
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6973592dd0>}

```

epoch: 41, training loss: 0.7754894639365375, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:47 model\_time: 0.8508 (0.8508) evaluator\_time:  
0.4560 (0.4560) time: 2.1442 data: 0.8190 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4983 (0.5272) evaluator\_time:  
0.2120 (0.2440) time: 0.8302 data: 0.0679 max mem: 5550

Test: Total time: 0:00:19 (0.8930 s / it)

Averaged stats: model\_time: 0.4983 (0.5272) evaluator\_time: 0.2120 (0.2440)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox

```

Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.386
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.629
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.431
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.376
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.621
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.344
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.509
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.551
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.316
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.542
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.755

```

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69611031d0>}

epoch: 42, training loss: 0.7545679155737162, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:44 model\_time: 0.4828 (0.4828) evaluator\_time:  
0.2548 (0.2548) time: 2.0312 data: 1.2832 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4634 (0.4868) evaluator\_time:  
0.2057 (0.2473) time: 0.7932 data: 0.0719 max mem: 5550

Test: Total time: 0:00:19 (0.8746 s / it)

Averaged stats: model\_time: 0.4634 (0.4868) evaluator\_time: 0.2057 (0.2473)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.432
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.377
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.346
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.510
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.556
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.321
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.552
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.751
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6972344210>}

```

epoch: 43, training loss: 0.7503184089437127, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:12 model\_time: 0.5184 (0.5184) evaluator\_time:  
0.2466 (0.2466) time: 3.2964 data: 2.5213 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4728 (0.4772) evaluator\_time:  
0.1963 (0.2314) time: 0.7634 data: 0.0630 max mem: 5550

Test: Total time: 0:00:19 (0.8964 s / it)

Averaged stats: model\_time: 0.4728 (0.4772) evaluator\_time: 0.1963 (0.2314)

Accumulating evaluation results...

DONE (t=0.57s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.389
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.433
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.094
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.377
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.344
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.507
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.553
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.323
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.546
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.750
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6a7518d050>}

```

epoch: 44, training loss: 0.7427291590720415, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:01 model\_time: 0.7865 (0.7865) evaluator\_time:  
0.3900 (0.3900) time: 2.8007 data: 1.6014 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4726 (0.4939) evaluator\_time:



```

0.1969 (0.2383)  time: 0.7545  data: 0.0638  max mem: 5550
Test: Total time: 0:00:19 (0.8811 s / it)
Averaged stats: model_time: 0.4726 (0.4939)  evaluator_time: 0.1969 (0.2383)
Accumulating evaluation results...
DONE (t=0.62s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.389
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.630
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.435
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.379
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.621
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.345
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.502
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.547
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.326
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.542
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.729
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69628d9c10>}

```

```

epoch: 45, training loss: 0.7431679600849748, lr: 0.0001
creating index...
index created!
Test: [ 0/22]  eta: 0:01:00  model_time: 0.8131 (0.8131)  evaluator_time:
0.4544 (0.4544)  time: 2.7395  data: 1.4530  max mem: 5550
Test: [21/22]  eta: 0:00:00  model_time: 0.4847 (0.5148)  evaluator_time:
0.1785 (0.2276)  time: 0.7535  data: 0.0636  max mem: 5550
Test: Total time: 0:00:19 (0.8941 s / it)
Averaged stats: model_time: 0.4847 (0.5148)  evaluator_time: 0.1785 (0.2276)
Accumulating evaluation results...
DONE (t=0.59s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.389
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.632
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.430
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.378
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.621
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.344
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.508
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.553
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.312
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.545
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.755
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69735f9210>}

```

```

epoch: 46, training loss: 0.7368227518163621, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:47 model_time: 0.7986 (0.7986) evaluator_time:
0.4565 (0.4565) time: 2.1642 data: 0.8934 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4657 (0.4950) evaluator_time:
0.1718 (0.2409) time: 0.7615 data: 0.0548 max mem: 5550
Test: Total time: 0:00:19 (0.8926 s / it)
Averaged stats: model_time: 0.4657 (0.4950) evaluator_time: 0.1718 (0.2409)
Accumulating evaluation results...
DONE (t=0.59s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.387
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.429
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.096
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.378
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.621
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.344
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.514
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.561
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.344
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.548
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.755
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69735fd250>}

```

```

epoch: 47, training loss: 0.7472575001884252, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:52 model_time: 0.7414 (0.7414) evaluator_time:
0.2899 (0.2899) time: 2.3824 data: 1.3301 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.5011 (0.5166) evaluator_time:
0.1709 (0.2262) time: 0.7798 data: 0.0673 max mem: 5550
Test: Total time: 0:00:19 (0.8862 s / it)
Averaged stats: model_time: 0.5011 (0.5166) evaluator_time: 0.1709 (0.2262)
Accumulating evaluation results...
DONE (t=0.59s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.630
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.431
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.380
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.621
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.345
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.509
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.556

```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.316
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.551
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.755
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6972231710>}
```

epoch: 48, training loss: 0.7292201286181808, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:11 model\_time: 0.7323 (0.7323) evaluator\_time:  
0.2957 (0.2957) time: 3.2667 data: 2.2211 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4699 (0.4887) evaluator\_time:  
0.1795 (0.2230) time: 0.7410 data: 0.0615 max mem: 5550

Test: Total time: 0:00:19 (0.8871 s / it)

Averaged stats: model\_time: 0.4699 (0.4887) evaluator\_time: 0.1795 (0.2230)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.389
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.628
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.430
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.376
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.346
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.503
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.551
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.318
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.552
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.734
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697206d110>}
```

epoch: 49, training loss: 0.7317307190969586, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:49 model\_time: 0.9977 (0.9977) evaluator\_time:  
0.3256 (0.3256) time: 2.2720 data: 0.9299 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4871 (0.5228) evaluator\_time:  
0.1829 (0.2217) time: 0.7857 data: 0.0571 max mem: 5550

Test: Total time: 0:00:19 (0.9001 s / it)

Averaged stats: model\_time: 0.4871 (0.5228) evaluator\_time: 0.1829 (0.2217)

Accumulating evaluation results...

DONE (t=0.63s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.628
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.432
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.380
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.348
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.506
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.552
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.321
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.551
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6961100ad0>}

```

epoch: 50, training loss: 0.7309305733069777, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:10 model\_time: 0.7068 (0.7068) evaluator\_time:  
0.1920 (0.1920) time: 3.2100 data: 2.2938 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4760 (0.5002) evaluator\_time:  
0.1670 (0.2077) time: 0.7337 data: 0.0629 max mem: 5550

Test: Total time: 0:00:19 (0.8875 s / it)

Averaged stats: model\_time: 0.4760 (0.5002) evaluator\_time: 0.1670 (0.2077)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.627
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.430
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.377
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.347
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.509
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.556
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.345
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.548
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6962ac0050>}

```

epoch: 51, training loss: 0.723961332347244, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:43 model\_time: 0.7600 (0.7600) evaluator\_time:  
0.3551 (0.3551) time: 1.9644 data: 0.8257 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4563 (0.5271) evaluator\_time:  
0.1912 (0.2378) time: 0.7945 data: 0.0682 max mem: 5550

Test: Total time: 0:00:19 (0.8886 s / it)

Averaged stats: model\_time: 0.4563 (0.5271) evaluator\_time: 0.1912 (0.2378)

Accumulating evaluation results...

DONE (t=0.59s).

IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area= all	maxDets=100 ]	= 0.391
Average Precision	(AP)	@[ IoU=0.50	area= all	maxDets=100 ]	= 0.628
Average Precision	(AP)	@[ IoU=0.75	area= all	maxDets=100 ]	= 0.430
Average Precision	(AP)	@[ IoU=0.50:0.95	area= small	maxDets=100 ]	= 0.096
Average Precision	(AP)	@[ IoU=0.50:0.95	area=medium	maxDets=100 ]	= 0.377
Average Precision	(AP)	@[ IoU=0.50:0.95	area= large	maxDets=100 ]	= 0.626
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all	maxDets= 1 ]	= 0.349
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all	maxDets= 10 ]	= 0.505
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all	maxDets=100 ]	= 0.552
Average Recall	(AR)	@[ IoU=0.50:0.95	area= small	maxDets=100 ]	= 0.323
Average Recall	(AR)	@[ IoU=0.50:0.95	area=medium	maxDets=100 ]	= 0.551
Average Recall	(AR)	@[ IoU=0.50:0.95	area= large	maxDets=100 ]	= 0.734

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697359ea50>}

epoch: 52, training loss: 0.7230595569126308, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:01 model\_time: 0.8615 (0.8615) evaluator\_time:  
0.2962 (0.2962) time: 2.8005 data: 1.6268 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4872 (0.5255) evaluator\_time:  
0.1812 (0.2262) time: 0.7697 data: 0.0618 max mem: 5550

Test: Total time: 0:00:19 (0.8993 s / it)

Averaged stats: model\_time: 0.4872 (0.5255) evaluator\_time: 0.1812 (0.2262)

Accumulating evaluation results...

DONE (t=0.59s).

IoU metric: bbox

Average Precision	(AP)	@[ IoU=0.50:0.95	area= all	maxDets=100 ]	= 0.390
Average Precision	(AP)	@[ IoU=0.50	area= all	maxDets=100 ]	= 0.629
Average Precision	(AP)	@[ IoU=0.75	area= all	maxDets=100 ]	= 0.431
Average Precision	(AP)	@[ IoU=0.50:0.95	area= small	maxDets=100 ]	= 0.096
Average Precision	(AP)	@[ IoU=0.50:0.95	area=medium	maxDets=100 ]	= 0.376
Average Precision	(AP)	@[ IoU=0.50:0.95	area= large	maxDets=100 ]	= 0.622
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all	maxDets= 1 ]	= 0.346
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all	maxDets= 10 ]	= 0.506
Average Recall	(AR)	@[ IoU=0.50:0.95	area= all	maxDets=100 ]	= 0.552
Average Recall	(AR)	@[ IoU=0.50:0.95	area= small	maxDets=100 ]	= 0.324
Average Recall	(AR)	@[ IoU=0.50:0.95	area=medium	maxDets=100 ]	= 0.550
Average Recall	(AR)	@[ IoU=0.50:0.95	area= large	maxDets=100 ]	= 0.735

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697233bb10>}

epoch: 53, training loss: 0.735398753080517, lr: 0.0001

creating index...

index created!

```

Test: [ 0/22] eta: 0:01:14 model_time: 0.6500 (0.6500) evaluator_time:
0.3308 (0.3308) time: 3.3874 data: 2.3971 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4674 (0.5087) evaluator_time:
0.1892 (0.2116) time: 0.7721 data: 0.0592 max mem: 5550
Test: Total time: 0:00:19 (0.8999 s / it)
Averaged stats: model_time: 0.4674 (0.5087) evaluator_time: 0.1892 (0.2116)
Accumulating evaluation results...
DONE (t=0.63s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.392
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.630
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.427
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.096
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.381
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.630
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.348
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.513
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.558
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.347
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.549
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.741
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69720478d0>}

```

```

epoch: 54, training loss: 0.7226636332925409, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:42 model_time: 0.7082 (0.7082) evaluator_time:
0.3100 (0.3100) time: 1.9358 data: 0.8988 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4527 (0.5163) evaluator_time:
0.1756 (0.2397) time: 0.8054 data: 0.0755 max mem: 5550
Test: Total time: 0:00:19 (0.8866 s / it)
Averaged stats: model_time: 0.4527 (0.5163) evaluator_time: 0.1756 (0.2397)
Accumulating evaluation results...
DONE (t=0.60s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.412
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.096
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.373
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.627
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.346
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.503
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.548
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.321
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.544
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735

```

```
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69735d9250>}
```

epoch: 55, training loss: 0.6978816972114146, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:51 model\_time: 0.8910 (0.8910) evaluator\_time:  
0.4527 (0.4527) time: 2.3183 data: 0.9550 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4664 (0.4974) evaluator\_time:  
0.2010 (0.2564) time: 0.7899 data: 0.0600 max mem: 5550

Test: Total time: 0:00:19 (0.8877 s / it)

Averaged stats: model\_time: 0.4664 (0.4974) evaluator\_time: 0.2010 (0.2564)

Accumulating evaluation results...

DONE (t=0.61s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.626

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.430

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.094

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.375

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.627

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.349

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.505

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.550

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.319

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.549

Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735

```
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69721ecc90>}
```

epoch: 56, training loss: 0.7311661839485168, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:40 model\_time: 0.6678 (0.6678) evaluator\_time:  
0.2169 (0.2169) time: 1.8515 data: 0.9523 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5091 (0.5281) evaluator\_time:  
0.1983 (0.2362) time: 0.8135 data: 0.0729 max mem: 5550

Test: Total time: 0:00:19 (0.8940 s / it)

Averaged stats: model\_time: 0.5091 (0.5281) evaluator\_time: 0.1983 (0.2362)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.391

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.629

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.428

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.097

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.376

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.629

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.349
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.512
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.556
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.345
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.549
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69720c4dd0>}

```

epoch: 57, training loss: 0.7259415695443749, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:02 model\_time: 0.9089 (0.9089) evaluator\_time:  
0.3746 (0.3746) time: 2.8269 data: 1.5265 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4837 (0.5387) evaluator\_time:  
0.1900 (0.2207) time: 0.8041 data: 0.0667 max mem: 5550

Test: Total time: 0:00:19 (0.9068 s / it)

Averaged stats: model\_time: 0.4837 (0.5387) evaluator\_time: 0.1900 (0.2207)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox

```

Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.389
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.627
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.421
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.098
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.373
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.625
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.345
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.505
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.551
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.323
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.547
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.736

```

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6961100190>}

epoch: 58, training loss: 0.7125398677308112, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:46 model\_time: 0.8457 (0.8457) evaluator\_time:  
0.3994 (0.3994) time: 2.1183 data: 0.8554 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4844 (0.5245) evaluator\_time:  
0.1657 (0.2258) time: 0.7868 data: 0.0702 max mem: 5550

Test: Total time: 0:00:19 (0.9017 s / it)

Averaged stats: model\_time: 0.4844 (0.5245) evaluator\_time: 0.1657 (0.2258)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox



```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.391
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.419
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.376
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.629
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.348
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.506
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.552
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.319
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.551
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6973583410>}

```

```

epoch: 59, training loss: 0.7150231753475964, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:52 model_time: 0.7241 (0.7241) evaluator_time:
0.3517 (0.3517) time: 2.3727 data: 1.2858 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4642 (0.5254) evaluator_time:
0.1773 (0.2215) time: 0.7839 data: 0.0635 max mem: 5550
Test: Total time: 0:00:19 (0.8813 s / it)
Averaged stats: model_time: 0.4642 (0.5254) evaluator_time: 0.1773 (0.2215)
Accumulating evaluation results...
DONE (t=0.59s).
IoU metric: bbox

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.389
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.628
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.432
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.374
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.346
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.512
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.557
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.349
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.548
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.736
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f697118add0>}

```

```

epoch: 60, training loss: 0.7208215785212815, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:43 model_time: 0.8371 (0.8371) evaluator_time:
0.3782 (0.3782) time: 1.9853 data: 0.7482 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4806 (0.5240) evaluator_time:

```

```

0.1942 (0.2373)  time: 0.7971  data: 0.0664  max mem: 5550
Test: Total time: 0:00:19 (0.8816 s / it)
Averaged stats: model_time: 0.4806 (0.5240)  evaluator_time: 0.1942 (0.2373)
Accumulating evaluation results...
DONE (t=0.58s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.391
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.631
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.420
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.094
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.374
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.629
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.348
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.505
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.551
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.326
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.545
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.740
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6973648150>}

```

```

epoch: 61, training loss: 0.7082522052805871, lr: 0.0001
creating index...
index created!
Test: [ 0/22]  eta: 0:00:42  model_time: 0.7316 (0.7316)  evaluator_time:
0.2982 (0.2982)  time: 1.9507  data: 0.9044  max mem: 5550
Test: [21/22]  eta: 0:00:00  model_time: 0.4529 (0.5255)  evaluator_time:
0.1648 (0.2227)  time: 0.7510  data: 0.0611  max mem: 5550
Test: Total time: 0:00:19 (0.8788 s / it)
Averaged stats: model_time: 0.4529 (0.5255)  evaluator_time: 0.1648 (0.2227)
Accumulating evaluation results...
DONE (t=0.57s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.389
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.629
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.416
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.094
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.371
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.629
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.348
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.507
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.552
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.329
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.545
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.741
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69710c0d90>}

```

```

epoch: 62, training loss: 0.7075089241843671, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:45 model_time: 0.7093 (0.7093) evaluator_time:
0.2765 (0.2765) time: 2.0874 data: 1.0901 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.5229 (0.5179) evaluator_time:
0.2105 (0.2444) time: 0.8327 data: 0.0769 max mem: 5550
Test: Total time: 0:00:19 (0.8994 s / it)
Averaged stats: model_time: 0.5229 (0.5179) evaluator_time: 0.2105 (0.2444)
Accumulating evaluation results...
DONE (t=0.61s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.387
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.630
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.417
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.094
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.373
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.348
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.505
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.550
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.331
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.545
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.734
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69735f9210>}

```

```

epoch: 63, training loss: 0.7019422589801252, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:52 model_time: 0.9637 (0.9637) evaluator_time:
0.5118 (0.5118) time: 2.3815 data: 0.8849 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4672 (0.5189) evaluator_time:
0.1902 (0.2318) time: 0.7766 data: 0.0689 max mem: 5550
Test: Total time: 0:00:19 (0.8880 s / it)
Averaged stats: model_time: 0.4672 (0.5189) evaluator_time: 0.1902 (0.2318)
Accumulating evaluation results...
DONE (t=0.58s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.631
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.417
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.097
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.373
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.628
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.346
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.513
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.558

```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.357
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.544
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.740
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6966144f10>}
```

epoch: 64, training loss: 0.7096419697627425, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:14 model\_time: 0.5039 (0.5039) evaluator\_time:  
0.1837 (0.1837) time: 3.3782 data: 2.6810 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4651 (0.4851) evaluator\_time:  
0.1855 (0.2183) time: 0.7650 data: 0.0652 max mem: 5550

Test: Total time: 0:00:19 (0.9017 s / it)

Averaged stats: model\_time: 0.4651 (0.4851) evaluator\_time: 0.1855 (0.2183)

Accumulating evaluation results...

DONE (t=0.58s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.630
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.429
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.376
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.348
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.507
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.553
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.323
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.549
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.742
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69735ba390>}
```

epoch: 65, training loss: 0.7018594392575324, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:49 model\_time: 0.9476 (0.9476) evaluator\_time:  
0.3778 (0.3778) time: 2.2355 data: 0.8929 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4575 (0.5308) evaluator\_time:  
0.1792 (0.2383) time: 0.7679 data: 0.0675 max mem: 5550

Test: Total time: 0:00:19 (0.8941 s / it)

Averaged stats: model\_time: 0.4575 (0.5308) evaluator\_time: 0.1792 (0.2383)

Accumulating evaluation results...

DONE (t=0.59s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.389
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.418
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.094
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.374
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.346
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.506
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.553
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.322
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.549
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.743
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6973543350>}

```

epoch: 66, training loss: 0.6957095521502197, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:44 model\_time: 0.6620 (0.6620) evaluator\_time: 0.3595 (0.3595) time: 2.0183 data: 0.9801 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4726 (0.5161) evaluator\_time: 0.2148 (0.2528) time: 0.8022 data: 0.0651 max mem: 5550

Test: Total time: 0:00:19 (0.8931 s / it)

Averaged stats: model\_time: 0.4726 (0.5161) evaluator\_time: 0.2148 (0.2528)

Accumulating evaluation results...

DONE (t=0.60s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.389
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.416
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.096
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.376
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.345
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.506
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.553
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.326
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.549
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697107af50>}

```

epoch: 67, training loss: 0.6778616174124181, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:48 model\_time: 0.8188 (0.8188) evaluator\_time: 0.4623 (0.4623) time: 2.1825 data: 0.8911 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4785 (0.5332) evaluator\_time: 0.1751 (0.2459) time: 0.7942 data: 0.0744 max mem: 5550

Test: Total time: 0:00:20 (0.9097 s / it)

Averaged stats: model\_time: 0.4785 (0.5332) evaluator\_time: 0.1751 (0.2459)

Accumulating evaluation results...

DONE (t=0.61s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.391
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.631
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.418
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.094
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.378
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.346
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.507
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.554
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.322
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.551
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.741
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f692c88e5d0>}
```

epoch: 68, training loss: 0.7089858367107809, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:44 model\_time: 0.8045 (0.8045) evaluator\_time:  
0.3349 (0.3349) time: 2.0202 data: 0.8699 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4783 (0.5177) evaluator\_time:  
0.2135 (0.2305) time: 0.8133 data: 0.0825 max mem: 5550

Test: Total time: 0:00:19 (0.8870 s / it)

Averaged stats: model\_time: 0.4783 (0.5177) evaluator\_time: 0.2135 (0.2305)

Accumulating evaluation results...

DONE (t=0.58s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.631
Average Precision (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.416
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.374
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.624
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.345
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.505
Average Recall    (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.551
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.321
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.547
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.741
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69735ae310>}
```

epoch: 69, training loss: 0.7045988992322236, lr: 0.0001

creating index...

index created!

```

Test: [ 0/22] eta: 0:00:40 model_time: 0.7351 (0.7351) evaluator_time:
0.3034 (0.3034) time: 1.8514 data: 0.7959 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4836 (0.5165) evaluator_time:
0.1926 (0.2428) time: 0.7967 data: 0.0644 max mem: 5550
Test: Total time: 0:00:19 (0.8762 s / it)
Averaged stats: model_time: 0.4836 (0.5165) evaluator_time: 0.1926 (0.2428)
Accumulating evaluation results...
DONE (t=0.60s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.631
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.419
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.097
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.377
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.621
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.341
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.511
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.558
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.353
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.549
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.734
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6971115f10>}

```

```

epoch: 70, training loss: 0.6852119578979909, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:48 model_time: 0.6545 (0.6545) evaluator_time:
0.2560 (0.2560) time: 2.2247 data: 1.3005 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4883 (0.5144) evaluator_time:
0.1839 (0.2334) time: 0.7780 data: 0.0647 max mem: 5550
Test: Total time: 0:00:19 (0.8851 s / it)
Averaged stats: model_time: 0.4883 (0.5144) evaluator_time: 0.1839 (0.2334)
Accumulating evaluation results...
DONE (t=0.60s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.391
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.631
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.418
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.379
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.341
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.506
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.554
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.322
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.552
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739

```

{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6961100a90>}

epoch: 71, training loss: 0.688182502053678, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:03 model\_time: 0.8013 (0.8013) evaluator\_time:  
0.3912 (0.3912) time: 2.8872 data: 1.6837 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4626 (0.5040) evaluator\_time:  
0.2217 (0.2205) time: 0.8029 data: 0.0879 max mem: 5550

Test: Total time: 0:00:19 (0.8982 s / it)

Averaged stats: model\_time: 0.4626 (0.5040) evaluator\_time: 0.2217 (0.2205)

Accumulating evaluation results...

DONE (t=0.59s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.391

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.631

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.420

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.097

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.379

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.341

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.507

Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.554

Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.326

Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.550

Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739

{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f6971231810>}

epoch: 72, training loss: 0.7010308583267033, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:42 model\_time: 0.8364 (0.8364) evaluator\_time:  
0.3282 (0.3282) time: 1.9485 data: 0.7693 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.5006 (0.5277) evaluator\_time:  
0.2113 (0.2264) time: 0.8079 data: 0.0538 max mem: 5550

Test: Total time: 0:00:19 (0.8914 s / it)

Averaged stats: model\_time: 0.5006 (0.5277) evaluator\_time: 0.2113 (0.2264)

Accumulating evaluation results...

DONE (t=0.61s).

IoU metric: bbox

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.389

Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.629

Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.416

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.098

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.378

Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622



```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.341
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.505
Average Recall      (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.552
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.325
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.550
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.734
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69735ca750>}

```

epoch: 73, training loss: 0.6900442542973906, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:19 model\_time: 0.9341 (0.9341) evaluator\_time:  
0.3925 (0.3925) time: 3.6047 data: 2.2603 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4562 (0.5127) evaluator\_time:  
0.1703 (0.2035) time: 0.7503 data: 0.0612 max mem: 5550

Test: Total time: 0:00:19 (0.8927 s / it)

Averaged stats: model\_time: 0.4562 (0.5127) evaluator\_time: 0.1703 (0.2035)

Accumulating evaluation results...

DONE (t=0.58s).

IoU metric: bbox

```

Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.391
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.629
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.418
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.098
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.379
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.624
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.341
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.506
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.554
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.324
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.555
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735

```

{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6962784890>}

epoch: 74, training loss: 0.6928675111848861, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:01:01 model\_time: 0.5542 (0.5542) evaluator\_time:  
0.2285 (0.2285) time: 2.7762 data: 1.9835 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4284 (0.4980) evaluator\_time:  
0.2220 (0.2262) time: 0.7910 data: 0.0629 max mem: 5550

Test: Total time: 0:00:19 (0.8904 s / it)

Averaged stats: model\_time: 0.4284 (0.4980) evaluator\_time: 0.2220 (0.2262)

Accumulating evaluation results...

DONE (t=0.58s).

IoU metric: bbox

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.391
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.630
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.418
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.379
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.624
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.342
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.507
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.554
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.328
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.552
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f69744ba450>}

```

```

epoch: 75, training loss: 0.6810111734084785, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:46 model_time: 0.7736 (0.7736) evaluator_time:
0.4158 (0.4158) time: 2.1112 data: 0.9043 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4695 (0.5266) evaluator_time:
0.2104 (0.2357) time: 0.8142 data: 0.0697 max mem: 5550
Test: Total time: 0:00:19 (0.8912 s / it)
Averaged stats: model_time: 0.4695 (0.5266) evaluator_time: 0.2104 (0.2357)
Accumulating evaluation results...
DONE (t=0.57s).
IoU metric: bbox

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.391
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.629
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.420
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.096
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.377
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.342
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.507
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.555
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.325
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.556
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.735
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6960f82b50>}

```

```

epoch: 76, training loss: 0.6898893732577562, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:01:05 model_time: 0.4910 (0.4910) evaluator_time:
0.2134 (0.2134) time: 2.9908 data: 2.2763 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4802 (0.5101) evaluator_time:

```

```

0.1751 (0.2150)  time: 0.7792  data: 0.0651  max mem: 5550
Test: Total time: 0:00:19 (0.9064 s / it)
Averaged stats: model_time: 0.4802 (0.5101)  evaluator_time: 0.1751 (0.2150)
Accumulating evaluation results...
DONE (t=0.59s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.392
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.628
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.430
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.091
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.383
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.349
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.510
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.553
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.329
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.548
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f697360a7d0>}

```

```

epoch: 77, training loss: 0.6910979349631816, lr: 0.0001
creating index...
index created!
Test: [ 0/22]  eta: 0:00:47  model_time: 0.5197 (0.5197)  evaluator_time:
0.2465 (0.2465)  time: 2.1583  data: 1.3822  max mem: 5550
Test: [21/22]  eta: 0:00:00  model_time: 0.4600 (0.5116)  evaluator_time:
0.1766 (0.2445)  time: 0.7957  data: 0.0609  max mem: 5550
Test: Total time: 0:00:19 (0.8928 s / it)
Averaged stats: model_time: 0.4600 (0.5116)  evaluator_time: 0.1766 (0.2445)
Accumulating evaluation results...
DONE (t=0.60s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.390
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.629
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.419
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.380
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.342
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.501
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.548
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.296
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.553
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6962b02750>}

```

```

epoch: 78, training loss: 0.6775415851734579, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:52 model_time: 0.8828 (0.8828) evaluator_time:
0.5119 (0.5119) time: 2.3862 data: 0.9743 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4948 (0.5277) evaluator_time:
0.1613 (0.2263) time: 0.7608 data: 0.0611 max mem: 5550
Test: Total time: 0:00:19 (0.8977 s / it)
Averaged stats: model_time: 0.4948 (0.5277) evaluator_time: 0.1613 (0.2263)
Accumulating evaluation results...
DONE (t=0.59s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.627
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.430
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.092
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.379
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.623
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.349
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.507
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.554
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.322
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.553
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.741
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f697360a7d0>}

```

```

epoch: 79, training loss: 0.6817199701908976, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:45 model_time: 0.6886 (0.6886) evaluator_time:
0.2486 (0.2486) time: 2.0613 data: 1.1068 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4933 (0.5205) evaluator_time:
0.1960 (0.2533) time: 0.8325 data: 0.0814 max mem: 5550
Test: Total time: 0:00:20 (0.9176 s / it)
Averaged stats: model_time: 0.4933 (0.5205) evaluator_time: 0.1960 (0.2533)
Accumulating evaluation results...
DONE (t=0.60s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.626
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.417
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.095
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.378
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.348
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.500
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.547

```

```
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.299
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.551
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.734
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6971086fd0>}
```

epoch: 80, training loss: 0.6823933951091021, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:50 model\_time: 0.7129 (0.7129) evaluator\_time:  
0.3560 (0.3560) time: 2.3032 data: 1.2165 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4845 (0.5181) evaluator\_time:  
0.1952 (0.2282) time: 0.7616 data: 0.0671 max mem: 5550

Test: Total time: 0:00:19 (0.8850 s / it)

Averaged stats: model\_time: 0.4845 (0.5181) evaluator\_time: 0.1952 (0.2282)

Accumulating evaluation results...

DONE (t=0.58s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.391
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.627
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.429
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.093
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.381
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.624
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.349
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.503
Average Recall    (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.549
Average Recall    (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.297
Average Recall    (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.554
Average Recall    (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.740
{'bbox': <pycocotools.cocoeval.COCOeval object at 0x7f6962ae0e10>}
```

epoch: 81, training loss: 0.6852330153342336, lr: 0.0001

creating index...

index created!

Test: [ 0/22] eta: 0:00:46 model\_time: 0.8172 (0.8172) evaluator\_time:  
0.3514 (0.3514) time: 2.1078 data: 0.9229 max mem: 5550

Test: [21/22] eta: 0:00:00 model\_time: 0.4814 (0.5234) evaluator\_time:  
0.1560 (0.2317) time: 0.7860 data: 0.0617 max mem: 5550

Test: Total time: 0:00:19 (0.8888 s / it)

Averaged stats: model\_time: 0.4814 (0.5234) evaluator\_time: 0.1560 (0.2317)

Accumulating evaluation results...

DONE (t=0.59s).

IoU metric: bbox

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.390
Average Precision (AP) @[ IoU=0.50      | area= all | maxDets=100 ] = 0.626
Average Precision (AP) @[ IoU=0.75      | area= all | maxDets=100 ] = 0.429
```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.091
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.379
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.624
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.356
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.499
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.547
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.295
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.553
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.736
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69722e1f50>}

```

```

epoch: 82, training loss: 0.6953576051164418, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:45 model_time: 0.8563 (0.8563) evaluator_time:
0.2746 (0.2746) time: 2.0554 data: 0.9117 max mem: 5550
Test: [21/22] eta: 0:00:00 model_time: 0.4290 (0.5511) evaluator_time:
0.1900 (0.2053) time: 0.8160 data: 0.0783 max mem: 5550
Test: Total time: 0:00:19 (0.8857 s / it)
Averaged stats: model_time: 0.4290 (0.5511) evaluator_time: 0.1900 (0.2053)
Accumulating evaluation results...
DONE (t=0.60s).

```

```

IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.388
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.625
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.427
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.091
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.380
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.622
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.347
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.508
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.557
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.327
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.556
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.739
{'bbox': <pycocotools.cocoeval.COCOEval object at 0x7f69610ada90>}

```

```

epoch: 83, training loss: 0.6889228262007236, lr: 0.0001
creating index...
index created!
Test: [ 0/22] eta: 0:00:44 model_time: 0.7438 (0.7438) evaluator_time:
0.4011 (0.4011) time: 2.0074 data: 0.8449 max mem: 5550

```

## 13 Inference

### 13.0.1 Load Model From Checkpoint

```
[ ]: SSD_MODEL = torchvision.models.detection.  
    ↳ssdlite320_mobilenet_v3_large(pretrained=False)  
  
checkPoint = torch.load("/content/gdrive/MyDrive/facemask_detection/save_model/  
    ↳facemaskDetectionSSD_320x320_38MAP.pth",  
                        map_location=getDevice())  
  
SSD_MODEL.load_state_dict(checkPoint['model_state_dict'])  
SSD_MODEL.eval()  
SSD_MODEL.to(getDevice())  
  
[ ]: SSD(  
    (backbone): SSDLiteFeatureExtractorMobileNet(  
        (features): Sequential(  
            (0): Sequential(  
                (0): ConvBNActivation(  
                    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),  
bias=False)  
                    (1): BatchNorm2d(16, eps=0.001, momentum=0.03, affine=True,  
track_running_stats=True)  
                    (2): Hardswish()  
                )  
                (1): InvertedResidual(  
                    (block): Sequential(  
                        (0): ConvBNActivation(  
                            (0): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1,  
1), groups=16, bias=False)  
                            (1): BatchNorm2d(16, eps=0.001, momentum=0.03, affine=True,  
track_running_stats=True)  
                            (2): ReLU(inplace=True)  
                        )  
                        (1): ConvBNActivation(  
                            (0): Conv2d(16, 16, kernel_size=(1, 1), stride=(1, 1), bias=False)  
                            (1): BatchNorm2d(16, eps=0.001, momentum=0.03, affine=True,  
track_running_stats=True)  
                            (2): Identity()  
                        )  
                    )  
                )  
                (2): InvertedResidual(  
                    (block): Sequential(  
                        (0): ConvBNActivation(  
                            (0): Conv2d(16, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
```

```

        (1): BatchNorm2d(64, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
    )
    (1): ConvBNActivation(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=64, bias=False)
        (1): BatchNorm2d(64, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
    )
    (2): ConvBNActivation(
        (0): Conv2d(64, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(24, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): Identity()
    )
)
)
(3): InvertedResidual(
  (block): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(24, 72, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(72, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
    )
    (1): ConvBNActivation(
      (0): Conv2d(72, 72, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=72, bias=False)
      (1): BatchNorm2d(72, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
    )
    (2): ConvBNActivation(
      (0): Conv2d(72, 24, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(24, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Identity()
    )
  )
)
(4): InvertedResidual(
  (block): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(24, 72, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(72, eps=0.001, momentum=0.03, affine=True,

```



```

track_running_stats=True)
    (2): ReLU(inplace=True)
    )
    (1): ConvBNActivation(
      (0): Conv2d(72, 72, kernel_size=(5, 5), stride=(2, 2), padding=(2,
2), groups=72, bias=False)
      (1): BatchNorm2d(72, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (fc1): Conv2d(72, 24, kernel_size=(1, 1), stride=(1, 1))
      (relu): ReLU(inplace=True)
      (fc2): Conv2d(24, 72, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): ConvBNActivation(
      (0): Conv2d(72, 40, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(40, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Identity()
    )
  )
)
(5): InvertedResidual(
  (block): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(40, 120, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(120, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
    )
    (1): ConvBNActivation(
      (0): Conv2d(120, 120, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=120, bias=False)
      (1): BatchNorm2d(120, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (fc1): Conv2d(120, 32, kernel_size=(1, 1), stride=(1, 1))
      (relu): ReLU(inplace=True)
      (fc2): Conv2d(32, 120, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): ConvBNActivation(
      (0): Conv2d(120, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)

```

```

        (1): BatchNorm2d(40, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): Identity()
    )
)
(6): InvertedResidual(
  (block): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(40, 120, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(120, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
    )
    (1): ConvBNActivation(
      (0): Conv2d(120, 120, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=120, bias=False)
      (1): BatchNorm2d(120, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
    )
    (2): SqueezeExcitation(
      (fc1): Conv2d(120, 32, kernel_size=(1, 1), stride=(1, 1))
      (relu): ReLU(inplace=True)
      (fc2): Conv2d(32, 120, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): ConvBNActivation(
      (0): Conv2d(120, 40, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(40, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Identity()
    )
  )
)
(7): InvertedResidual(
  (block): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(40, 240, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(240, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Hardswish()
    )
    (1): ConvBNActivation(
      (0): Conv2d(240, 240, kernel_size=(3, 3), stride=(2, 2),

```

```

padding=(1, 1), groups=240, bias=False)
    (1): BatchNorm2d(240, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()
    )
    (2): ConvBNActivation(
    (0): Conv2d(240, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(80, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Identity()
    )
    )
    (8): InvertedResidual(
    (block): Sequential(
    (0): ConvBNActivation(
    (0): Conv2d(80, 200, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(200, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()
    )
    (1): ConvBNActivation(
    (0): Conv2d(200, 200, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=200, bias=False)
    (1): BatchNorm2d(200, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()
    )
    (2): ConvBNActivation(
    (0): Conv2d(200, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(80, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Identity()
    )
    )
    (9): InvertedResidual(
    (block): Sequential(
    (0): ConvBNActivation(
    (0): Conv2d(80, 184, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(184, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()

```

```

        )
        (1): ConvBNActivation(
          (0): Conv2d(184, 184, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=184, bias=False)
          (1): BatchNorm2d(184, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): Hardswish()
        )
        (2): ConvBNActivation(
          (0): Conv2d(184, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(80, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): Identity()
        )
      )
    )
    (10): InvertedResidual(
      (block): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(80, 184, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(184, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): Hardswish()
        )
        (1): ConvBNActivation(
          (0): Conv2d(184, 184, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=184, bias=False)
          (1): BatchNorm2d(184, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): Hardswish()
        )
        (2): ConvBNActivation(
          (0): Conv2d(184, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
          (1): BatchNorm2d(80, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): Identity()
        )
      )
    )
    (11): InvertedResidual(
      (block): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)

```

```

        (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): Hardswish()
    )
    (1): ConvBNActivation(
        (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=480, bias=False)
        (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): Hardswish()
    )
    (2): SqueezeExcitation(
        (fc1): Conv2d(480, 120, kernel_size=(1, 1), stride=(1, 1))
        (relu): ReLU(inplace=True)
        (fc2): Conv2d(120, 480, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): ConvBNActivation(
        (0): Conv2d(480, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): Identity()
    )
)
)
(12): InvertedResidual(
    (block): Sequential(
        (0): ConvBNActivation(
            (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1),
bias=False)
            (1): BatchNorm2d(672, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
            (2): Hardswish()
        )
        (1): ConvBNActivation(
            (0): Conv2d(672, 672, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), groups=672, bias=False)
            (1): BatchNorm2d(672, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
            (2): Hardswish()
        )
        (2): SqueezeExcitation(
            (fc1): Conv2d(672, 168, kernel_size=(1, 1), stride=(1, 1))
            (relu): ReLU(inplace=True)
            (fc2): Conv2d(168, 672, kernel_size=(1, 1), stride=(1, 1))
        )
        (3): ConvBNActivation(

```

```

        (0): Conv2d(672, 112, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (1): BatchNorm2d(112, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): Identity()
    )
)
)
(13): ConvBNActivation(
    (0): Conv2d(112, 672, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(672, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()
)
)
(1): Sequential(
  (0): Sequential(
    (1): ConvBNActivation(
      (0): Conv2d(672, 672, kernel_size=(5, 5), stride=(2, 2), padding=(2,
2), groups=672, bias=False)
      (1): BatchNorm2d(672, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Hardswish()
    )
    (2): SqueezeExcitation(
      (fc1): Conv2d(672, 168, kernel_size=(1, 1), stride=(1, 1))
      (relu): ReLU(inplace=True)
      (fc2): Conv2d(168, 672, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): ConvBNActivation(
      (0): Conv2d(672, 80, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(80, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Identity()
    )
  )
)
(1): InvertedResidual(
  (block): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Hardswish()
    )
    (1): ConvBNActivation(
      (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),

```

```

padding=(2, 2), groups=480, bias=False)
    (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()
    )
    (2): SqueezeExcitation(
    (fc1): Conv2d(480, 120, kernel_size=(1, 1), stride=(1, 1))
    (relu): ReLU(inplace=True)
    (fc2): Conv2d(120, 480, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): ConvBNActivation(
    (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(80, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Identity()
    )
    )
    (2): InvertedResidual(
    (block): Sequential(
    (0): ConvBNActivation(
    (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()
    )
    (1): ConvBNActivation(
    (0): Conv2d(480, 480, kernel_size=(5, 5), stride=(1, 1),
padding=(2, 2), groups=480, bias=False)
    (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Hardswish()
    )
    (2): SqueezeExcitation(
    (fc1): Conv2d(480, 120, kernel_size=(1, 1), stride=(1, 1))
    (relu): ReLU(inplace=True)
    (fc2): Conv2d(120, 480, kernel_size=(1, 1), stride=(1, 1))
    )
    (3): ConvBNActivation(
    (0): Conv2d(480, 80, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (1): BatchNorm2d(80, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): Identity()
    )
    )

```

```

    )
    )
    (3): ConvBNActivation(
      (0): Conv2d(80, 480, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): Hardswish()
    )
  )
  )
  (extra): ModuleList(
    (0): Sequential(
      (0): ConvBNActivation(
        (0): Conv2d(480, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU6(inplace=True)
      )
      (1): ConvBNActivation(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=256, bias=False)
        (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU6(inplace=True)
      )
      (2): ConvBNActivation(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(512, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU6(inplace=True)
      )
    )
  )
  (1): Sequential(
    (0): ConvBNActivation(
      (0): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(128, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU6(inplace=True)
    )
    (1): ConvBNActivation(
      (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=128, bias=False)
      (1): BatchNorm2d(128, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
      (2): ReLU6(inplace=True)
    )
    (2): ConvBNActivation(

```



```

        (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU6(inplace=True)
    )
)
(2): Sequential(
  (0): ConvBNActivation(
    (0): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): ReLU6(inplace=True)
  )
  (1): ConvBNActivation(
    (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), groups=128, bias=False)
    (1): BatchNorm2d(128, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): ReLU6(inplace=True)
  )
  (2): ConvBNActivation(
    (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): ReLU6(inplace=True)
  )
)
(3): Sequential(
  (0): ConvBNActivation(
    (0): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): ReLU6(inplace=True)
  )
  (1): ConvBNActivation(
    (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
groups=64, bias=False)
    (1): BatchNorm2d(64, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): ReLU6(inplace=True)
  )
  (2): ConvBNActivation(
    (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(128, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
    (2): ReLU6(inplace=True)
  )
)

```

```

    )
    )
    )
    (anchor_generator): DefaultBoxGenerator(aspect_ratios=[[2, 3], [2, 3], [2, 3],
[2, 3], [2, 3], [2, 3]], clip=True, scales=[0.2, 0.35, 0.5, 0.65, 0.8, 0.95,
1.0], steps=None)
    (head): SSDLiteHead(
      (classification_head): SSDLiteClassificationHead(
        (module_list): ModuleList(
          (0): Sequential(
            (0): ConvBNActivation(
              (0): Conv2d(672, 672, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=672, bias=False)
              (1): BatchNorm2d(672, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
              (2): ReLU6(inplace=True)
            )
            (1): Conv2d(672, 546, kernel_size=(1, 1), stride=(1, 1))
          )
          (1): Sequential(
            (0): ConvBNActivation(
              (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=480, bias=False)
              (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
              (2): ReLU6(inplace=True)
            )
            (1): Conv2d(480, 546, kernel_size=(1, 1), stride=(1, 1))
          )
          (2): Sequential(
            (0): ConvBNActivation(
              (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=512, bias=False)
              (1): BatchNorm2d(512, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
              (2): ReLU6(inplace=True)
            )
            (1): Conv2d(512, 546, kernel_size=(1, 1), stride=(1, 1))
          )
          (3): Sequential(
            (0): ConvBNActivation(
              (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=256, bias=False)
              (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
              (2): ReLU6(inplace=True)
            )
          )
        )
      )
    )
  )
)

```

```

        (1): Conv2d(256, 546, kernel_size=(1, 1), stride=(1, 1))
    )
    (4): Sequential(
      (0): ConvBNActivation(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=256, bias=False)
        (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU6(inplace=True)
      )
      (1): Conv2d(256, 546, kernel_size=(1, 1), stride=(1, 1))
    )
    (5): Sequential(
      (0): ConvBNActivation(
        (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=128, bias=False)
        (1): BatchNorm2d(128, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
        (2): ReLU6(inplace=True)
      )
      (1): Conv2d(128, 546, kernel_size=(1, 1), stride=(1, 1))
    )
  )
  (regression_head): SSDLiteRegressionHead(
    (module_list): ModuleList(
      (0): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(672, 672, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=672, bias=False)
          (1): BatchNorm2d(672, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(672, 24, kernel_size=(1, 1), stride=(1, 1))
      )
      (1): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(480, 480, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=480, bias=False)
          (1): BatchNorm2d(480, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(480, 24, kernel_size=(1, 1), stride=(1, 1))
      )
      (2): Sequential(

```

```

        (0): ConvBNActivation(
          (0): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=512, bias=False)
          (1): BatchNorm2d(512, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(512, 24, kernel_size=(1, 1), stride=(1, 1))
      )
      (3): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=256, bias=False)
          (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(256, 24, kernel_size=(1, 1), stride=(1, 1))
      )
      (4): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=256, bias=False)
          (1): BatchNorm2d(256, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(256, 24, kernel_size=(1, 1), stride=(1, 1))
      )
      (5): Sequential(
        (0): ConvBNActivation(
          (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), groups=128, bias=False)
          (1): BatchNorm2d(128, eps=0.001, momentum=0.03, affine=True,
track_running_stats=True)
          (2): ReLU6(inplace=True)
        )
        (1): Conv2d(128, 24, kernel_size=(1, 1), stride=(1, 1))
      )
    )
  )
)
(transform): GeneralizedRCNNTransform(
  Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5])
  Resize(min_size=(320,), max_size=320, mode='bilinear')
)
)

```

### 13.0.2 Show Output

```
[ ]: device = getDevice()
device
```

```
[ ]: device(type='cuda')
```

```
[ ]: from matplotlib import pyplot as plt
import matplotlib.patches as patches

def show_output_with_bbox(img, bboxes, labels, scores):

    print(img.shape)

    img_height = img.shape[0]
    img_width = img.shape[1]

    fig, ax = plt.subplots(figsize=(8,8))

    ax.imshow(img)

    for box, class_name, score in zip(bboxes, labels, scores):

        if score > 0.1:

            class_name = class_dict[class_name]
            xmin, ymin, xmax, ymax = box
            xmin, ymin, width, height = [x for x in [xmin, ymin, (xmax-xmin),
↪(ymax-ymin)]]
            rect = patches.Rectangle((xmin, ymin), width, height, linewidth=1,
↪edgecolor='r', facecolor='none')
            ax.add_patch(rect)
            ax.text(xmin, ymin, f"{class_name, np.round(score, 2)}", color='r',
↪fontsize='large')

            # {'xx-small', 'x-small', 'small', 'medium', 'large', 'x-large',
↪'xx-large'}

    plt.show()

#####
```

```

[ ]: test_images = VALID_DF['filename'].unique().tolist()
train_images = TRAIN_DF['filename'].unique().tolist()

for index, image in enumerate(test_images):

    if index == 10:
        break

    class_dict = {
        0: "background",
        1: "with_mask",
        2: "without_mask",
        3: "mask_wearred_incorrect"
    }

    img = cv2.imread(test_images[index])
    img = cv2.resize(img, (320, 320))
    img = cv2.cvtColor(img, cv2.COLOR_BGRA2RGB).astype(np.float32)
    img = img/255.0

    SSD_MODEL.eval()
    img_tensor_float = torch.tensor(img, dtype=torch.float32).permute(2,0,1).
    ↳unsqueeze(dim=0).to(device)
    out = SSD_MODEL(img_tensor_float)

    keepIndex = torchvision.ops.nms(out[0]['boxes'], out[0]['scores'],
    ↳iou_threshold=0.2)

    bboxTensors = out[0]['boxes'][keepIndex]
    labelsTensors = out[0]['labels'][keepIndex]
    scoresTensors = out[0]['scores'][keepIndex]

    boxes = bboxTensors.detach().cpu().numpy().astype(np.int)
    labels = labelsTensors.detach().cpu().numpy()
    scores = scoresTensors.detach().cpu().numpy()

    show_output_with_bbox(img, boxes, labels, scores)

```

### 13.0.3 From Single Image

```
[ ]: from matplotlib import pyplot as plt
import matplotlib.patches as patches

def show_output_with_bbox(img, bboxes, labels, scores):

    print(img.shape)

    img_height = img.shape[0]
    img_width = img.shape[1]

    fig, ax = plt.subplots(figsize=(8,8))

    ax.imshow(img)

    for box, class_name, score in zip(bboxes, labels, scores):

        if score > 0.2:

            class_name = class_dict[class_name]
            xmin, ymin, xmax, ymax = box
            xmin, ymin, width, height = [x for x in [xmin, ymin, (xmax-xmin),
→(ymax-ymin)]]
            rect = patches.Rectangle((xmin, ymin), width, height, linewidth=1,
→edgecolor='r', facecolor='none')
            ax.add_patch(rect)
            ax.text(xmin, ymin, f"{class_name, np.round(score, 2)}", color='r',
→fontsize='large')

            # {'xx-small', 'x-small', 'small', 'medium', 'large', 'x-large',
→'xx-large'}

    plt.show()

#####

[ ]: imagePath = "/content/Screenshot from 2021-10-14 21-25-39.png"

img = cv2.imread(imagePath)
img = cv2.resize(img, (320, 320))
img = cv2.cvtColor(img, cv2.COLOR_BGRA2RGB).astype(np.float32)
img = img/255.0
```

```

SSD_MODEL.eval()
img_tensor_float = torch.tensor(img, dtype=torch.float32).permute(2,0,1).
    ↳unsqueeze(dim=0).to(device)
out = SSD_MODEL(img_tensor_float)

keepIndex = torchvision.ops.nms(out[0]['boxes'], out[0]['scores'],
    ↳iou_threshold=0.1)

bboxTensors = out[0]['boxes'][keepIndex]
labelsTensors = out[0]['labels'][keepIndex]
scoresTensors = out[0]['scores'][keepIndex]

boxes = bboxTensors.detach().cpu().numpy().astype(np.int)
labels = labelsTensors.detach().cpu().numpy()
scores = scoresTensors.detach().cpu().numpy()

show_output_with_bbox(img, boxes, labels, scores)

```

(320, 320, 3)





## 14 Deployed Heroku API

```
[ ]: import requests

file = "/content/IMG_20211011_022805_1.jpg"
resp = requests.post("https://facemask-detection-api.herokuapp.com/predict",
    ↪files={'filename': open(file, 'rb')})

[ ]: resp.content

[ ]: b'{"result":{"boxes": [[100,151,234,306], [116,195,219,304], [68,86,255,291], [133,158,226,273], [114,183,157,231], [203,141,227,166], [218,142,238,164], [113,160,163,2
```

21], [127, 182, 189, 219], [129, 157, 185, 217], [135, 162, 204, 201], [205, 130, 226, 151], [111, 290, 205, 320], [131, 188, 163, 214], [149, 159, 198, 214], [139, 196, 193, 235], [192, 170, 234, 259], [89, 286, 117, 320], [137, 177, 224, 225], [239, 123, 258, 146], [110, 149, 203, 246], [155, 207, 183, 242], [75, 108, 153, 248], [239, 140, 257, 158], [135, 197, 173, 230], [198, 159, 232, 221], [196, 125, 213, 147], [221, 120, 239, 141], [112, 250, 211, 301], [139, 208, 166, 237], [238, 112, 263, 140], [197, 139, 216, 164], [237, 130, 260, 153], [168, 165, 227, 248], [78, 153, 171, 279], [200, 142, 221, 162], [215, 129, 238, 152], [144, 177, 198, 237], [189, 197, 233, 231], [105, 6, 241, 146], [141, 293, 227, 320], [210, 140, 245, 166], [146, 152, 207, 185], [157, 255, 238, 302], [239, 152, 256, 172], [218, 99, 243, 131], [21, 122, 130, 235], [201, 147, 223, 166], [163, 165, 221, 204], [84, 296, 104, 320], [121, 206, 213, 261], [45, 70, 169, 213], [45, 284, 101, 320], [192, 184, 235, 214], [205, 113, 226, 137], [196, 142, 224, 163], [62, 292, 82, 320], [138, 190, 222, 240], [213, 135, 242, 155], [116, 196, 152, 227], [204, 141, 235, 165], [131, 168, 175, 197], [221, 128, 258, 151], [212, 146, 240, 169], [10, 140, 132, 284], [126, 149, 196, 196], [170, 198, 200, 233], [220, 152, 238, 172], [177, 295, 242, 320], [36, 178, 283, 314], [155, 187, 185, 218], [23, 132, 193, 309], [237, 89, 264, 130], [194, 112, 213, 137], [146, 213, 194, 261], [194, 139, 228, 163], [72, 36, 204, 174], [209, 147, 243, 190], [256, 120, 281, 156], [169, 106, 193, 140], [188, 133, 219, 153], [102, 174, 144, 233], [157, 283, 233, 309], [35, 0, 284, 181], [238, 102, 259, 133], [84, 18, 192, 140], [191, 286, 242, 309], [143, 7, 279, 150], [117, 279, 210, 311], [202, 170, 231, 199], [23, 85, 132, 202], [179, 191, 229, 240], [256, 134, 281, 172], [114, 188, 174, 242], [93, 292, 160, 320], [175, 140, 205, 168], [77, 77, 202, 224], [229, 119, 276, 163], [103, 210, 154, 250], [0, 50, 22, 107], [254, 103, 281, 143], [212, 83, 256, 133], [99, 183, 143, 275], [240, 168, 256, 193], [0, 161, 21, 215], [3, 82, 155, 320], [211, 164, 237, 207], [26, 284, 83, 320], [121, 223, 213, 279], [168, 120, 238, 265], [0, 49, 105, 171], [161, 113, 180, 140], [57, 298, 103, 320], [161, 162, 210, 222], [59, 273, 84, 320], [160, 98, 180, 128], [194, 137, 224, 156], [193, 102, 211, 127], [44, 290, 66, 320], [0, 129, 21, 183], [0, 145, 21, 199], [178, 209, 242, 252], [125, 153, 177, 184], [219, 19, 249, 66], [0, 191, 22, 248], [118, 114, 320, 320], [217, 297, 271, 320], [198, 147, 229, 180], [0, 112, 21, 167], [234, 18, 266, 66], [0, 179, 21, 230], [45, 8, 174, 153], [84, 294, 132, 320], [0, 96, 21, 152], [124, 185, 189, 251], [209, 10, 258, 73], [172, 95, 193, 126], [240, 156, 260, 185], [227, 99, 277, 149], [14, 35, 147, 182], [137, 209, 169, 253], [240, 1, 262, 32], [118, 50, 232, 167], [207, 102, 224, 125], [0, 0, 168, 252], [219, 156, 238, 181], [107, 20, 136, 66], [29, 298, 92, 320], [123, 20, 153, 66], [54, 48, 161, 169], [0, 33, 22, 91], [72, 296, 116, 320], [0, 10, 29, 76], [257, 148, 280, 186], [214, 162, 247, 212], [0, 215, 24, 284], [139, 20, 168, 66], [203, 20, 232, 65], [0, 0, 34, 55], [255, 0, 278, 33], [155, 20, 184, 66], [224, 228, 270, 313], [222, 92, 242, 122], [224, 1, 245, 32], [187, 20, 216, 65], [0, 70, 28, 141], [200, 144, 229, 189], [171, 20, 200, 66], [221, 64, 244, 104], [183, 87, 224, 132], [119, 185, 149, 222], [219, 111, 250, 133], [193, 10, 242, 72], [217, 220, 308, 316], [97, 10, 145, 73], [199, 246, 235, 305], [8, 7, 90, 150], [113, 10, 161, 73], [27, 289, 53, 320], [129, 10, 177, 72], [249, 54, 320, 168], [95, 194, 156, 295], [11, 285, 36, 320], [145, 10, 193, 72], [208, 1, 229, 32], [151, 0, 320, 256], [177, 10, 225, 72], [161, 10, 209, 72], [176, 1, 197, 32], [160, 1, 181, 32], [144, 1, 165, 32], [192, 1, 213, 32], [197, 86, 239, 131], [128, 1, 149, 32], [205, 67, 228, 104], [112, 1, 133, 32], [238, 0, 320, 91], [210, 19, 317, 137], [227, 78, 276, 136], [203, 285, 231, 320], [208, 149, 235, 171], [236, 297, 297, 320], [208, 91, 236, 116], [102, 278, 170, 311], [223, 8, 275, 74], [157, 63, 181, 106], [235, 59, 261, 104], [44, 34, 286, 312], [189, 67, 212, 105], [76, 0, 213, 89], [8, 156, 40, 220], [164, 82, 206, 132], [0, 0, 29, 34], [172, 65, 196, 105], [90, 19, 121, 66], [297, 36, 320, 116], [6, 183, 41, 255], [0, 6, 20, 43], [124, 60, 149, 105], [140, 60, 165, 105], [0, 0, 84, 91], [137, 291, 169, 320], [149, 290, 182, 320], [12, 3, 35, 39], [228, 131, 273, 178], [61, 62, 85, 105], [217, 167, 234, 196], [108, 61, 133, 104], [224, 91, 253, 116], [298, 157, 320, 239], [247, 12, 283, 70], [

[illegible]

32519945129752,0.001514149596914649,0.0015070142690092325,0.0015022751176729798,  
0.0015014008386060596,0.0014999363338574767,0.0014962218701839447,0.001496217329  
9863935,0.0014923999551683664,0.0014787133550271392,0.0014750001719221473,0.0014  
697472797706723,0.0014632644597440958,0.001461434061639011,0.0014608936617150903  
,0.0014431339222937822,0.0014232895337045193,0.0014157508267089725,0.00141174893  
36058497,0.001409144140779972,0.0014082142151892185,0.001404500100761652,0.00140  
11674793437123,0.0014011637540534139,0.0013938621850684285,0.0013840317260473967  
,0.0013821367174386978,0.0013742017326876521,0.0013721405994147062,0.00136986514  
5534277,0.0013648574240505695,0.0013627833686769009,0.0013619204983115196,0.0013  
617892982438207,0.0013569580623880029,0.0013567578280344605,0.001356201712042093  
3,0.001352174673229456,0.0013458242174237967,0.001334638916887343,0.001332846237  
34653,0.0013309161877259612,0.0013284091837704182,0.0013235552469268441,0.001313  
957036472857,0.0013130586594343185,0.0013127595884725451,0.0013122461969032884,0  
.0013020304031670094,0.0012994048884138465,0.0012987260706722736,0.0012928050709  
88834,0.001286111306399107,0.001284601166844368,0.0012826028978452086,0.00127838  
10961991549,0.0012761630350723863,0.001273144269362092,0.0012644161470234394,0.0  
012639377964660525,0.0012577680172398686,0.0012569597456604242,0.001249959226697  
6833,0.0012498770374804735,0.0012492816895246506,0.001245302613824606,0.00124503  
46257537603,0.0012399195693433285,0.001238647266291082,0.001238569151610136,0.00  
12384398141875863,0.0012375334044918418,0.001237453194335103,0.00123608321882784  
37,0.0012332525802776217,0.001231589587405324,0.001228789333254099,0.00122769270  
09224892,0.0012267718557268381,0.0012253267923370004,0.0012192405993118882,0.001  
218406017869711,0.0012091726530343294,0.0012067040661349893,0.00120108004193753,  
0.001198226003907621,0.001192128867842257,0.0011858935467898846,0.00118454534094  
78068,0.0011827201815322042,0.0011815419420599937,0.0011806224938482046,0.001177  
8734624385834,0.001175625016912818,0.0011726247612386942,0.0011710436083376408,0  
.0011654995614662766,0.0011623450554907322,0.001159784966148436,0.00115737796295  
43424,0.0011538844555616379,0.0011533374199643731,0.0011476489016786218,0.001147  
0596073195338,0.0011392548913136125,0.0011305513326078653,0.0011298366589471698,  
0.0011274936841800809,0.0011259699240326881,0.0011199618456885219,0.001119256718  
0842161,0.0011175676481798291,0.0011106460588052869,0.0011071154149249196,0.0011  
065710568800569,0.0011043440317735076,0.0011026815045624971,0.001099870656616985  
8,0.0010975177865475416,0.0010936460457742214,0.0010910271666944027,0.0010889532  
277360559,0.0010823864722624421,0.0010737490374594927,0.0010713861556723714,0.00  
10675836820155382,0.0010652048513293266,0.0010643661953508854,0.0010634816717356  
443,0.0010626617586240172,0.001054805121384561,0.0010538556380197406,0.001048661  
65202111,0.0010480402270331979,0.0010408080415800214,0.0010378758888691664,0.001  
0377306025475264,0.001037541194818914,0.001037459704093635,0.0010355329141020775  
,0.0010343752801418304,0.0010325741022825241,0.0010296597611159086,0.00102814624  
55168366,0.0010207975283265114,0.0010206286096945405,0.0010200011311098933,0.001  
0193875059485435,0.0010099031496793032,0.0010091960430145264}], "success": true}\n

[ ]: